
Tokenizer-Attention Coupling: How BPE Merge Decisions Permanently Shape Transformer Internal Organization

dayna@blackwell-systems.com · DOI:
10.5281/zenodo.20925910

Dayna Blackwell, Blackwell Systems

2026-07-02

Abstract

BPE tokenizers merge delimiter characters with adjacent content, destroying structural boundaries before the model sees them. We show that this has a larger consequence than previously understood: it permanently constrains the attention capacity of every head in the transformer, not just the small number of heads that visibly specialize on delimiters.

In controlled experiments across two architectures (GPT-NeoX, Llama), two model scales (410M, 1.3B), and three domains (structured data, programming languages, molecular notation), we establish the following:

1. **Stranded attention heads.** When a standard BPE model is given input with clean delimiter boundaries (using the model's own vocabulary, no retraining), all 384 attention heads show 4x more delimiter attention (14% to 55%). The model has the circuitry for structural processing. The tokenizer prevents it from working.
2. **The frustration gap is immediate and permanent.** The gap between the model's structural capacity and what the tokenizer allows appears by training step 5,000 and does not change across 35,000 additional steps. Training cannot compensate for tokenizer-imposed stranding.
3. **Merge barriers: a 16-character fix.** Forbidding 16 delimiter characters from BPE merges eliminates stranding. Models trained with merge barriers develop 50-161 concentrated structural anchoring heads (13-21% of capacity, scaling with model size), achieve 3-738x better structured data perplexity, 1.5x on code, and 2.2x on molecular notation, with zero natural language cost.
4. **Causal proof.** An 18-phase ablation protocol proves these heads are necessary (+59% to +420% degradation when removed), sufficient (70 heads outperform all 384), and format-general (8/9 unseen formats transfer).
5. **Scaling strengthens the mechanism.** At 1.3B parameters, 21% of attention heads specialize on delimiters (up from 13% at 410M), and ablation degradation increases 20x within the Llama architecture. Standard BPE models develop counterproductive delimiter heads at scale: removing them improves comprehension by 57%.
6. **Domain independence.** Per-token loss advantage in code maps directly to BPE merge rates (parentheses at 100% merge rate show 1.5x advantage; characters below 65% show reversal). In chemistry, merge barriers produce 67% more valid generated molecules and 2.5x fewer valence errors.
7. **Tokenizer-attention coupling.** We propose this term for the general principle: BPE merge decisions made during tokenizer training causally constrain which attention heads develop and how much capacity is available for content processing. Confirmed across three domains, supported by independent validation from Su et al. (2025) and unintentional confirmation from Jindal & Ju (2026).

The finding reframes the role of the tokenizer: it is not a preprocessing step but an architectural constraint that permanently shapes the model's internal organization. The fix is trivial (16 lines of tokenizer configuration), but the damage it addresses affects every model trained with standard BPE on structured input.

1. Introduction

1.1 The Problem

BPE tokenization was designed for natural language compression. It works by iteratively merging the most frequent adjacent character pairs into single vocabulary entries (Sennrich et al., 2016). This is effective: common words become single tokens, sequence lengths shrink, and models train faster.

But BPE has a blind spot. Structural delimiter characters (quotes, braces, parentheses, colons, pipes) co-occur with adjacent content at high frequency. BPE merges them. The double quote in "name" becomes part of a single token "name in 30% of production tokenizers. The opening parenthesis in process (data) merges with the function name at 100% in o200k_base. Every structural boundary that BPE merges is a boundary the model can no longer see at the token level.

This is not a minor inefficiency. Attention heads operate on token positions. If a structural boundary does not exist as a distinct token, no attention head can attend to it. The model must spend capacity rediscovering boundaries the tokenizer destroyed, a process we call boundary recovery.

Prior work observed symptoms of this problem. Su et al. (2025) found that delimiter choice swings structured data accuracy by $\pm 23\%$ but could not explain why. Zhang et al. (2026) showed that models count delimiters to locate table cells. Sandoval-Segura et al. (2025) identified dormant attention heads that correlate with input characteristics. None connected these observations to the tokenizer.

1.2 What We Found

We ran controlled experiments where two identical models (same architecture, same data, same hyperparameters) were trained with different tokenizers: one with BPE merge barriers (16 delimiter characters forbidden from merging) and one with standard BPE. The tokenizer was the only variable.

The merge-barrier model develops concentrated structural attention heads that lock onto clean delimiter tokens and use them to navigate structured input. We call this structural anchoring. The standard model's heads attempt the same specialization but fail because the delimiter tokens they need are fused with content. We call these stranded heads.

The stranding is more severe than initially apparent. When we fed the standard model input with clean delimiter boundaries (using its own vocabulary, no retraining, no weight changes), every single attention head (384/384 at 410M, 768/768 at 1.3B) showed 4x more delimiter attention. The model is not incapable of structural processing. The tokenizer is preventing it.

The stranding is also permanent. Measuring across 8 checkpoints from step 5,000 to step 40,000, the gap between what the model can do (55% delimiter attention with clean input) and what the tokenizer allows (14% with normal input) is flat. 35,000 steps of training, approximately 2.3 billion tokens

of structured data, do not narrow the gap by a single percentage point. We call this the frustration gap: 41 percentage points of structural attention capacity that the model builds, the gradient pushes toward, and the tokenizer permanently blocks.

The mechanism generalizes across domains. On programming languages, per-token loss advantage maps directly to BPE merge rates (parentheses at 100% merge rate show 1.5x advantage). On SMILES molecular notation, merge barriers produce 67% more valid generated molecules. Effect magnitude correlates with delimiter density across domains.

1.3 Contributions

This paper makes seven contributions:

1. We identify **stranded attention heads**, a new failure mode where heads are active but unproductive because corrupted token boundaries defeat their structural specialization. We prove stranding is caused by the tokenizer (384/384 heads wake up under clean tokenization with frozen weights), is immediate (present by step 5,000), and is permanent (unchanged through step 40,000). We show stranding is a stable state distinct from dormancy (no P0 drift across training). At 1.3B scale, standard BPE develops 124 counterproductive delimiter heads whose removal improves comprehension by 57% (Section 8.5).
2. We prove the **tokenizer-to-attention causal link**: BPE merge decisions during tokenizer training determine which attention heads develop during model training. Controlled experiments where the tokenizer is the only variable produce 50-161 structural anchoring heads (merge barriers) versus 3-124 stranded or counterproductive heads (standard BPE) (Sections 5-7).
3. We demonstrate **merge barriers**, a 16-character constraint on BPE, that eliminates stranding and produces 3-738x better structured data perplexity with zero natural language cost (Section 4).
4. We provide **causal ablation proof** (18-phase protocol) that delimiter heads are necessary (+59% to +420% degradation), sufficient (70 heads outperform 384), and format-general (8/9 unseen formats transfer) (Section 7).
5. We show the mechanism is **architecture-independent** (GPT-NeoX, Llama), **scale-dependent** (strengthens from 410M to 1.3B: 13% to 21% specialization, 20x more causal importance), and **domain-independent** (structured data 46-738x, code 1.5x, chemistry 2.2x) (Sections 8, 10, 11).
6. We answer open questions in two published papers: Su et al.'s delimiter performance ranking maps 1:1 to BPE merge rates; Jindal & Ju's character-level SMILES tokenizer was an implicit merge barrier (Section 11.2).
7. We propose **tokenizer-attention coupling** as the general principle: the tokenizer is not a pre-processing step but an architectural constraint that permanently shapes the model's internal organization. Confirmed across three domains (Section 11.3).

1.4 Scope and Limitations

All experiments use models at 410M and 1.3B parameters. We do not claim that the effect persists at 70B or 405B, though the scaling trend (strengthening from 410M to 1.3B) suggests it does. All PPL comparisons are single-run (no multi-seed replication), following standard practice in mechanistic interpretability where training costs make replication impractical (Olsson et al., 2022; Voita et al., 2019). We rely on internal consistency (per-token correlations, cross-format transfer, causal ablation) rather than confidence intervals. The PPL-to-comprehension gap remains open: we show better perplexity, not better downstream task accuracy on a standard benchmark.

We use Graph Compact Format (GCF), a header-factored wire format with pipe delimiters, as the comparison format. GCF was selected because its grammar characters have near-zero merge rates across all 43 tested tokenizers.

2. Background

2.1 BPE Tokenizers

Modern LLMs use Byte-Pair Encoding (BPE) tokenizers (Sennrich et al., 2016) trained on large text corpora, with SentencePiece (Kudo and Richardson, 2018) as the dominant implementation. Alternative subword methods exist, including Unigram language models (Kudo, 2018), but BPE remains the standard for GPT-family and Llama-family models. BPE builds a vocabulary by iteratively merging the most frequent byte sequences. The result is a fixed lookup table mapping strings to integer IDs. At inference time, the tokenizer greedily matches the longest vocabulary entry at each position. If "name exists as entry #32586, the tokenizer always selects it as one token. This is deterministic: a dictionary lookup, not a context-dependent decision.

2.2 Grammar vs. Payload

Any structured format contains grammar symbols (delimiters defining structure) and payload content (data values). In JSON, the grammar symbols are " , : , , { , } , [,] . Grammar symbols repeat on every row; payload content varies.

When grammar symbols merge with payload during BPE training, the resulting token conflates structural markup with semantic content. The token "name (GPT-4 #32586) encodes both “opening quote” (grammar: a field boundary starts here) and “name” (payload: this is the field called name). The model must decompose these two meanings from within a single embedding rather than reading the boundary from a dedicated grammar token. Each row of a JSON array contains the same "name": pattern, producing the same merged token. The ambiguity compounds linearly with data size: 10 rows produce 10 merged boundaries to decompose, 500 rows produce 3,000. At scale, grammar tokens dominate the input sequence (Section 3.4), and the merged boundaries make it increasingly difficult for the model to distinguish structure from content.

2.3 Structural Equivalence

Two models achieve structural equivalence when they see field boundaries at the same token positions. They may tokenize values differently (semantic variance, which is harmless), but they agree on where structure is. For example, GPT-4 and Claude both tokenize the value "pending" differently (GPT-4: 2 tokens, Claude: 3 tokens), but this does not affect comprehension because both models know where the value starts and ends.

When models disagree on where fields start and end, they are parsing different structures from the same input. GPT-4 sees ["value"] [":"] (boundary inside first token), while Claude sees [""] [value] [":"] (boundary at token edge). These models receive different structural signals from the same data, which produces model-dependent comprehension failures. A format achieves structural equivalence when all tokenizers place field boundaries at the same token positions. Our analysis shows GCF achieves 99.5% structural equivalence across all 43 tokenizers (Section 3.5); JSON achieves 7.5%.

2.4 Related Work: Attention Head Specialization

General head specialization. Voita et al. (2019) identified positional, syntactic, and rare-word heads in BERT. Michel et al. (2019) demonstrated that most heads can be pruned without significant loss. Olsson et al. (2022) proved that “induction heads” are causally responsible for in-context learning, building on the mathematical framework for transformer circuits (Elhage et al., 2021). Conmy et al. (2023) developed automated methods for discovering such circuits. Wang et al. (2025) tracked how heads differentiate into specialized roles during training. These studies are descriptive: they observe what heads do in existing models. None address what training conditions cause heads to specialize, and none study structured data formats.

Delimiter-related observations. Clark et al. (2019) observed that BERT heads attend heavily to [SEP] tokens, but [SEP] is a synthetic marker injected by the tokenizer, not a structural delimiter in data. Pande et al. (2021) classified heads into “syntactic, local, block, and delimiter” roles, using “delimiter” to mean [SEP]/[CLS] attention. Neither performed causal ablation to prove these heads are functionally necessary. The attention sink literature (Xiao et al., 2024) notes that models dump probability mass on delimiters and low-information tokens, but this is a training artifact (the model needs somewhere to put unused attention), not functional specialization for parsing.

Structured data attention. Zhang et al. (2026) found that LLMs use “an ordinal mechanism that counts discrete delimiters to resolve coordinates” in linearized tables, identifying specific multiplex heads. Wang et al. (2026) described three-phase attention in table understanding (scan, localize, amplify). Both study tables, not general structured formats, and neither connects attention behavior to tokenizer design. Su et al. (2025) showed delimiter choice affects aggregate attention by $\pm 23\%$ in ICL prompts but identified no specific heads and performed no causal intervention. Hasani et al. (2026) found dedicated counting heads that attend to commas but studied counting, not format comprehension. None of these works proposed or tested a tokenizer intervention.

Formal language parsing. Work on Dyck languages (Yao et al., 2021; Sharma et al., 2026) proves transformers can match brackets, and Jindal & Ju (2026) found a single head specializing on bracket

matching in SMILES molecular notation. These are formal-language studies on synthetic or narrow-domain data, not empirical analysis of how production models handle real serialized formats.

Code attention. Wan et al. (2022) showed attention in CodeBERT “aligns strongly with syntax structure” but reports aggregate correlation, not individual head identification on specific delimiter tokens. Katzy et al. (2026) mined attention patterns in Java code without characterizing delimiter-specific heads.

Dormant heads and attention sinks. Xiao et al. (2024) identified attention sinks (initial tokens absorbing disproportionate attention) as a training artifact. Sandoval-Segura et al. (2025) formalized “dormant heads” (heads dominated by attention sinks whose output can be zeroed without accuracy loss), finding 4-16% of heads dormant across six production models. Critically, they show dormancy correlates with input characteristics: structured text (code, math, citations) produces fewer dormant heads, while prose produces more. They explain 29.5% of variance via “alphabet density” (an inverse proxy for delimiter density) but explicitly state “additional research is needed to determine whether a single measure can accurately predict the fraction of dormant heads.”

Their dormant heads and our stranded heads represent two failure modes on a spectrum determined by tokenizer-attention coupling. Dormancy occurs when there is nothing structural to do: heads safely collapse into attention sinks (prose, high alphabet density, few delimiters). Stranding occurs when structural work IS needed but the tokenizer makes it impossible: heads remain active but unproductive, attending diffusely to merged tokens that hide boundary signals. Our stranded head experiment (Section 8.5) provides direct evidence for this spectrum: 16 heads in Model B show focused attention on clean delimiter input (GCF) but diffuse attention on merged input (JSON), demonstrating the intermediate state between productive anchoring and full dormancy. The distinction matters for inference: dormant heads can be safely zeroed (Sandoval-Segura et al. confirm this); stranded heads cannot, because the model is still attempting structural processing through them.

The gap we fill. No prior work has: (1) identified attention heads specializing on structural delimiters across multiple serialized data formats (JSON, CSV, GCF, YAML), (2) performed causal ablation proving these heads are necessary for structured data comprehension, (3) connected tokenizer merge behavior to the emergence of delimiter-specialized heads, or (4) demonstrated a tokenizer intervention that eliminates the need for these heads. We fill all four gaps with controlled experiments across two architectures.

Table 1: Contribution scope and prior-art boundaries.

Claim	Scope	Nearest prior work	What is new
Delimiter head identification	Multi-format (JSON, CSV, GCF, YAML)	Clark/Pande: BERT [SEP] only	Real structural delimiters
Causal ablation	+59% degradation, 18 phases	Su: correlational. Zhang: no ablation	Necessity proof for structured data

Claim	Scope	Nearest prior work	What is new
Tokenizer-attention link	Merge behavior determines heads	Chen: training data, not tokenizer	Vocabulary design causes specialization
Intervention eliminating heads	Merge barriers free capacity	No prior intervention	First fix demonstrated
Architecture independence	MHA and GQA both replicate	No comparable study	KV-group ablation methodology

3. The Problem: 43 Tokenizers, Universal Merging

3.1 Tokenizers Tested

We tested 43 tokenizers from 20 providers: OpenAI (cl100k, o200k, GPT-2), Anthropic (Claude), Meta (LLaMA 2/3/3.1, CodeLlama, TinyLlama), Google (Gemma 2/3, T5), Mistral (7B v0.1/v0.3, Nemo, Mixtral, Codestral), Alibaba (Qwen 2/2.5/3, QwQ), DeepSeek (V2/V3/R1), Microsoft (Phi-2/3/4), TII (Falcon), 01.AI (Yi), BigCode (StarCoder2), NVIDIA (Nemotron), AI21 (Jamba), Stability (StableLM), EleutherAI (Pythia), Snowflake (Arctic), AllenAI (OLMo), and Alibaba AIDC (Marco-o1). Vocabulary sizes range from 32K to 262K.

3.2 Field Boundary Merge Rates




Figure 1: Figure 1: Delimiter merge rates across 43 tokenizers




Figure 2: Figure 2: Field merge rates across 43 tokenizers

The most common JSON field names merge with the opening quote on 30% of tokenizers:

Field Pattern	Merge Rate	Affected Families
"id":, "name":, "time":, "title":	30.2% (13/43)	GPT-4, GPT-4o, LLaMA 3.x, Qwen, Phi-4, StableLM, Mistral Nemo
"type":, "value":, "url":, "text":	27.9% (12/43)	Same minus Mistral Nemo

According to Web Data Commons (University of Mannheim, 2024), name is the #1 most common JSON property on the web (3.5 billion occurrences). It merges on 30% of tokenizers.

On real evaluation data across 43 tokenizers (check counts vary by format due to different field counts):

Format	Merge Rate	Total Checks
GCF (pipe)	0.47%	29,025
JSON (quote)	8.17%	1,935
TOON (tab)	32.91%	860

The merge mechanism The variance has a specific cause: BPE merging absorbs the opening quote into the field name. When GPT-4’s tokenizer encounters "value": "pending", it produces:

["value"] [":"] [pending] ["] (4 tokens)

Claude’s tokenizer produces:

["] [value] [":"] [pending] ["] (5 tokens)

The structural boundary (where the field name starts) is at a different token position. On GPT-4, the opening quote is fused with content. On Claude, it is separate. The model must learn to decompose the merged token "value into “opening quote followed by a field name” rather than treating it as a single semantic unit.

By contrast, GCF’s pipe delimiter is always its own token on all 43 tokenizers:

value pending	-> [value][][pending]	ALL 43 tokenizers
name Alice	-> [name][][Alice]	ALL 43 tokenizers
orderId ORD-001	-> [orderId][][ORD][-][001]	ALL 43 tokenizers

Specific token IDs These are actual dictionary entries with specific IDs, not hypothetical merges:

Pattern	GPT-4	GPT-4o	LLaMA 3	Qwen 2.5	DeepSeek	Mistral	Claude	Gemma
"id	#29800	#60094	#29800	#28700	–	#117579	–	–
"name	#32586	#74800	#32586	#31486	–	#117753	–	–
"type	#45570	#91290	#45570	#44470	–	–	–	–
"value	#64407	#180654	#64407	#63307	–	–	–	–
"url	#61360	#124415	#61360	#60260	–	–	–	–

Cross-verified: encoding "name": "Alice" with GPT-4’s tokenizer confirms token #32586 is selected.

Maximum tokenization variance Searching across 840 JSON field+value patterns, the maximum variance case is "userName": "req_xyz789", which produces 7 distinct tokenizations across 8 models:

GPT-4, LLaMA:	["] [userName] [" : "] [req] [_ xyz] [789] ["]
GPT-4o:	[" user] [Name] [" : "] [req] [_ xyz] [789] ["]
Claude:	["] [userName] [" : "] [req] [_] [xyz] [789] ["]
Qwen 2.5:	["] [userName] [" : "] [req] [_ xyz] [7] [8] [9] ["]
DeepSeek V3:	["] [user] [Name] [" : "] [req] [_] [xyz] [789] ["]
Gemma 2:	["] [userName] [" : "] [req] [_] [xyz] [7] [8] [9] ["]
Mistral Nemo:	["] [user] [Name] [" : "] [req] [_ x] [yz] [7] [8] [9] ["]

A complete JSON object { "orderId": "ORD-001", "value": "shipped" } produces 4 different token counts (12, 13, 14, 15) depending on the model. The same data is literally a different length on different model families, affecting attention patterns, positional encodings, and context budget.

3.3 Adversarial Surface

We decoded every entry in all 43 vocabularies and classified entries where delimiter characters fuse with alphabetic content:

Delimiter	Unique Mergeable Words	Used by
\ (pipe)	24	GCF
" (quote)	193	JSON
: (colon)	232	JSON
, (comma)	282	JSON
\ t (tab)	1,238	TOON

JSON's total adversarial surface across all 7 grammar characters: **1,939 words**. That is 81x the pipe's 24. GPT-4 has 1,173 tab+letter vocabulary entries. TOON chose the delimiter with the largest adversarial surface of any common separator character.




Figure 3: Figure 3: Vocabulary merge entries by tokenizer

Multi-grammar vocabulary entries Some vocabulary entries contain multiple JSON grammar symbols fused together:

Token	Grammar Operations	Present In
" :	close string, start key-value, open string	43/43 tokenizers
" , "	close string, separate elements, open string	43/43 tokenizers
{ "	open object, open string	43/43 tokenizers
" : { "	close string, key-value, open object, open string	43/43 tokenizers
" } ,	close string, close object, separate	43/43 tokenizers
} , { "	close object, separate, open object, open string	33/43 tokenizers

The token " : { " packs four structural operations into one integer. The model receives one token where there should be four grammar decisions.

Pipe merge entries The pipe has 24 merged entries across all 43 vocabularies, but exclusively with programming keywords from type union syntax: | null, | string, | max, | min, | required. The entries | name, | id, | type, | value, | status, | title do not exist in any tested vocabulary.

The pipe merges with type-system keywords, not with field names. This is a dictionary fact, not a statistical claim.




Figure 4: ASCII adversarial surface across all 94 printable characters

ASCII character safety ranking We scanned all 94 printable ASCII characters (codes 33-126) across all 43 vocabularies:

Tier	Characters	Mergeable words
Safe (0 words)	0-9	0 (digits never merge with adjacent text)
Low risk (1-10)	` ~	5-8
Medium (11-50)	^ \ !] # % ?	17-50 (pipe at 24)
High (51-100)	& ; + } {	57-97
Very high (101+)	" : , [(- . _ letters	117-11,891

Digits are the only perfectly safe characters, but cannot serve as delimiters (they appear in payload data). Backtick (5 words) and tilde (8 words) have smaller surfaces than pipe (24), but both have practical drawbacks (backtick conflicts with markdown and template literals; tilde with paths and bitwise negation). The pipe was selected because its 24 mergeable words are all TypeScript union keywords, none of which appear as data field names, and it provides superior visual readability as a column separator.

3.4 Token Overhead




Figure 5: Figure 5: JSON overhead scaling: $O(n)$ per row vs $O(1)$ for header-factored formats

At 500 rows, JSON’s token distribution:

Category	% of Tokens
Repeated field names	52.4%
Structural characters	28.6%
Actual data values	19.0%

81% overhead. Only 19% carries information.

GCF for the same 500-row data:

Category	% of Tokens
Section header (declared once)	0.4%
Delimiter characters	1.9%
Actual data values	97.7%

GCF's overhead is 2.3%. JSON's is 81%. The difference is structural: GCF declares field names once in a header (`## orders [500]{orderId,customer,status,total}`) rather than repeating them on every row.




Figure 6: Figure 6: Savings stability bands across 43 tokenizers

Grammar swap experiment To confirm that token savings are a structural property (header factor-ing) and not an artifact of specific delimiter choices, we replaced all GCF delimiters with 4 alternative sets (all drawn from the non-merging character set) and re-measured savings across 5 payload types, 4 sizes, and 8 tokenizers (800 measurements). The spread across all delimiter sets was 0.4 percentage points. The savings come from eliminating repeated field names, not from using a particular delimiter character.

3.5 Structural Equivalence Proof



We measured structural equivalence by testing whether each grammar character tokenizes as its own token (isolated) or fuses with adjacent content (merged) across all 43 tokenizers.

GCF grammar maintains 99.5% isolation: @ is isolated on 43/43 tokenizers (100%), < on 43/43 (100%), | on 42.66/43 (99.2%; the 0.8% comes from | null and | string entries on tokenizers trained heavily on TypeScript). Every GCF field boundary is at a token edge on virtually every production tokenizer. Two different models processing the same GCF payload see the same structural boundaries.

JSON grammar fuses into multi-operation tokens on 43/43 tokenizers. Of all vocabulary entries containing the quote character, 92.5% encode multiple grammar operations (e.g., " : " encodes close-string + key-value-separator + open-string as a single token). The quote is isolated on only 3 tokenizers (Claude, Gemma 2, Gemma 3). On the remaining 40, field boundaries are inside merged tokens, producing the model-dependent tokenization variance documented in Section 3.2.

The structural equivalence gap (99.5% vs 7.5%) is not a property of specific tokenizers or training data; it is a property of which characters the format uses as delimiters. Any format using pipe, @, and < as grammar inherits the 99.5% isolation. Any format using quote, colon, and comma inherits the 7.5%. This is why merge barriers (Section 4) fix the problem for all formats simultaneously: isolating the 16 most-used delimiter characters brings every format closer to 100% structural equivalence.

3.6 Irrecoverability

The problem cannot be fixed for existing models:

1. Vocabulary is frozen post-training.
2. All weights depend on the vocabulary (token #32586 has learned embeddings in every layer).
3. Tokenization occurs before the transformer processes input.
4. Changing the tokenizer requires retraining the model from scratch.

No amount of fine-tuning, RLHF, or prompt engineering can change the fact that "name is a single token in GPT-4's dictionary.

3.7 Attention Mechanism on Pre-Existing Models

Before presenting our fix, we establish the transformer-level mechanism by which tokenizer merging causes comprehension failure, using pre-existing models (Pythia 410M, Gemma 2B).

Entropy crossover Attention entropy measures how spread out the model's attention is. High entropy means diffuse attention (looking everywhere, finding nothing). Low entropy means focused attention (knows where to look).

At small scale (5-20 orders), JSON entropy is lower than GCF. The model has been trained on billions of JSON examples and has efficient attention patterns. At 50 orders, the crossover: JSON entropy exceeds GCF by 13%. The model's learned JSON parsing breaks down as thousands of identical token IDs compete for attention.

Grammar attention collapse We classified every token as grammar or payload and measured attention allocation:

At small scale, JSON attention splits roughly 30% grammar / 68% payload. The model attends to structural tokens to understand the format. At 50 orders, JSON grammar attention collapses from 30% to 8.6%. The model stops attending to structural tokens. It distributes attention uniformly across content, unable to distinguish structure from data.

This is the mechanism behind comprehension failure. It is measurable, reproducible, and directly caused by the tokenizer producing merged boundary tokens that become indistinguishable at scale.

Ildiz et al. (2024) proved mathematically that self-attention weights tokens proportionally to their frequency in the input sequence. Their Context-Conditioned Markov Chain (CCMC) formulation shows that $P(\text{next_token} = j \mid X)$ includes m_j (the count of token j) in the numerator. When structural tokens like "name" : account for 80% of occurrences in a 500-row JSON array, they dominate the attention budget by count, leaving proportionally less for data values. The paper analyzes single-layer models; our comprehension data confirms the effect persists in production multi-layer architectures at 500+ rows.

Comprehension correlation The tokenization analysis connects to observed outcomes from independent comprehension evaluations across 10 frontier models (Blackwell, 2026):

- JSON accuracy at stress scale (500 symbols, 200 edges): **54.1%**
- GCF accuracy at stress scale (500 symbols, 200 edges): **91.2%**
- GCF accuracy on standard workloads (500 nested orders): **100%** on every frontier model

Error magnitude confirms the mechanism: GCF errors are small (off by 1-2, precision errors). JSON errors are large (off by 50-140, comprehension failures). The model did not slightly misread a number; it could not find the answer.

3.8 External Validation: Delimiter Performance Correlates with Merge Rate

Su et al. (2025) showed that changing a single delimiter character in ICL prompts swings model accuracy by $\pm 23\%$ across Llama, Qwen, and Gemma families. They measured that high-performing delimiters “steer attention towards key tokens” but stated the mechanism “warrants further research.”

We measure BPE merge rates for their delimiter set. A delimiter “merges” when the tokenizer absorbs it into the adjacent word as a single token (e.g., space + “name” → one token name). The correlation with their independently measured performance is exact:

Delimiter	BPE Merge Rate (o200k)	Boundary Status	Su et al. Performance
Space	100%	Always merges with following word	Worst performer

Delimiter	BPE Merge Rate (o200k)	Boundary Status	Su et al. Performance
Tab \t	90%	Merges on 9/10 common words	Poor
Comma ,	80%	Merges on 8/10 common words	Poor
Hyphen –	90%	Merges on 9/10 common words	Poor
Colon :	60%	Merges on 6/10 common words	Mixed
Semicolon ;	0%	Never merges	Good
Newline \n	0%	Never merges	Best performer
Pipe	0%	Never merges	Good
Hash #	0%	Never merges	Good

Every delimiter with 0% merge rate is a top performer in their evaluation. Every delimiter with 60-100% merge rate is a poor performer. The ranking is not model-dependent, not scale-dependent, and not task-dependent: it is a deterministic property of the tokenizer’s vocabulary.

This confirms our mechanism from an independent experiment: when a delimiter merges with adjacent content, the structural boundary disappears from the token stream. The model must allocate attention to recover it. When a delimiter stays as its own token (0% merge rate), the boundary is explicit and no recovery is needed. Su et al. measured the downstream effect ($\pm 23\%$ accuracy). We identify the root cause (BPE merging) and demonstrate the fix (merge barriers, which guarantee 0% merge rate for all 16 structural delimiters).

Reproducible: merge rates computed via tiktoken (o200k_base) on 10 common English words. Full 43-tokenizer analysis in `eval/hf-tokenizer-analysis.py`.

4. The Fix: Merge Barriers

4.1 Concept

Standard BPE (Sennrich et al., 2016) iteratively merges the most frequent adjacent byte pairs in the training corpus. When the corpus contains JSON, the pair " + n (from "name", "null", etc.) eventually crosses the frequency threshold and merges into a single vocabulary entry. Once merged, the tokenizer always selects it: "name becomes token #32586 in GPT-4. The merge is permanent and deterministic.

Merge barriers add a single constraint: 16 delimiter characters are forbidden from participating in any merge operation. The BPE algorithm itself is unchanged; the merge candidate list simply excludes pairs containing a barrier character. During tokenizer training, the pair " + n is never considered, regardless of its frequency. The barrier is enforced at the pre-tokenization stage (Kudo and Richardson, 2018), before BPE merging begins.

The result: every barrier character is always its own token. "name can never become a single token because " cannot merge with n. The model always sees explicit structural boundaries. This is the same guarantee that byte-level models like ByT5 (Xue et al., 2022) achieve by eliminating subword tokenization entirely, but without sacrificing BPE’s compression efficiency on natural language.

The tradeoff is quantifiable. Merge barriers prevent delimiters from merging into adjacent content, but the freed content bytes can form different (sometimes more efficient) merges. The net effect depends on delimiter density: on GCF, the merge-barrier tokenizer produces 14-19% fewer total tokens (113 vs 131 at 5 records, 1,855 vs 2,288 at 100 records); on natural language, ~3% more. The comprehension benefit (3-46x lower perplexity on structured data) far exceeds any compression cost.

4.2 Barrier Characters

16 characters selected for maximum structural coverage across structured data formats and code:

#	Character	Purpose
1	\	Field delimiter (GCF, shell pipes, markdown tables)
2	@	Symbol IDs (GCF), email, decorators
3	<	Edge direction (GCF), HTML/XML tags, comparisons
4	>	Edge direction (GCF), HTML/XML tags, comparisons
5	"	String delimiter (JSON, YAML)
6	'	String delimiter (YAML, code)
7	:	Key-value separator (JSON, YAML, Python)
8	,	Field separator (JSON, CSV, function arguments)
9	;	Statement terminator (code, CSV alternate)
10	\t	Column delimiter (TSV, TOON, indentation)
11	{	Open object (JSON, YAML), open block (code)
12	}	Close object (JSON, YAML), close block (code)
13	[	Open array (JSON), indexing
14	]	Close array (JSON), indexing
15	(	Open group (function calls, expressions)
16	)	Close group (function calls, expressions)

This is not a format-specific fix. It fixes JSON, YAML, CSV, TOON, GCF, and any format that uses delimiter characters.

4.3 Implementation

In HuggingFace tokenizers (the SentencePiece-compatible library used by most open-source LLMs), merge barriers are implemented as pre-tokenization rules that isolate barrier characters before BPE merging begins. Each barrier character is wrapped in a `Split(char, behavior="isolated")` pre-tokenizer, which segments the input so that the character appears alone in its own segment. When BPE then processes each segment independently, the barrier character has no adjacent bytes to merge with. The full configuration is in Appendix A.

This approach requires no modifications to the BPE algorithm, the training loop, or the model architecture. The constraint is entirely in the pre-tokenization stage, making it compatible with any BPE-based training pipeline. The 16 `Split` rules compose in a `Sequence` pre-tokenizer, followed by the standard `ByteLevel` pre-tokenizer that handles the remaining text. Training time is unchanged; the only difference is the input segmentation.

The resulting tokenizer has 65,539 vocabulary entries (3 more than the 65,536 standard-BPE control; most barrier characters already had byte-level entries, only 3 required new dedicated entries). It produces slightly more tokens on structured data (each delimiter is its own token rather than merging into adjacent content) but the same or fewer tokens on natural language, where the barrier characters appear infrequently.

4.4 Validation

We validated the trained tokenizer (structok-64k) through three independent checks:

1. **Exhaustive vocabulary scan.** Decoded every entry in the 65,539-entry vocabulary and verified that no entry contains a barrier character fused with alphabetic content. Result: **zero** merged delimiter entries. Compare to GPT-4's 117 quote+letter entries and 1,173 tab+letter entries.
2. **Adversarial surface measurement.** Applied the same adversarial surface analysis used on all 43 production tokenizers (Section 3.3). Result: **zero** adversarial surface. No barrier character appears inside any vocabulary entry alongside a letter. The attack surface that spans 1,939 words on JSON's grammar and 1,238 words on tab is eliminated entirely.
3. **Boundary isolation checks.** Tested 521 field+value patterns across all 16 barrier characters, verifying that each barrier character tokenizes as its own token in every context. Result: 521/521 passed. The isolation is deterministic: because barrier characters cannot participate in merges, they are always their own tokens regardless of surrounding content.

5. Controlled Experiments

5.1 Design

We trained four models in two controlled pairs, each differing only in the tokenizer. The experimental design isolates a single variable: within each pair, the architecture, corpus, hyperparameters, random seed, and hardware are identical. The only difference is whether 16 delimiter characters can participate in BPE merges.

Pair	Architecture	Steps	Context	Params
Run-002	GPT-NeoX 410M (24 layers, 16 heads, 1024 hidden)	20,000	2,048	436M
Run-003	Llama 410M (RoPE, GQA 4:1 [16 query/4 KV per layer], SwiGLU, RMSNorm)	40,000	2,048	405M

Both pairs used the same corpus (4.5 GB; major components: 33% FineWeb, 13% code, 14% JSON, 8% GCF, 1% YAML/CSV, 3% Wikipedia; remaining 28% additional FineWeb). Same hyperparameters within each pair.

Run-002 (NeoX) detail

	Model A (merge barriers)	Model B (standard BPE)
Architecture	GPT-NeoX 410M (436M params)	GPT-NeoX 410M (436M params)
Tokenizer	structok-64k (65,539 vocab, 16 barriers)	standard-64k (65,536 vocab, no barriers)
Training data	Same corpus (4.5 GB)	Same corpus (4.5 GB)
Pre-tokenized	1,258,728,671 tokens	1,269,271,190 tokens
Steps	20,000	20,000
Batch size	32 effective (8 x 4 GPUs)	32 effective
Learning rate	3e-4 flat	3e-4 flat
Hardware	4x A100 PCIE 40GB	4x A100 PCIE 40GB
Training	DDP with NCCL, gradient checkpointing, fp16	Same
Final overall PPL	19.4	19.5

Note: “Final overall PPL” is the perplexity observed during training (averaged across batches at step 20,000). The checkpoint stores the final batch loss, which differs from the running average.

We chose two architectures to test architecture independence. GPT-NeoX (Black et al., 2022) uses learned position embeddings, full multi-head attention (separate Q/K/V per head), GELU activation, and LayerNorm. Llama (Touvron et al., 2023) uses RoPE (Su et al., 2021), Grouped Query Attention (Ainslie et al., 2023; extending Shazeer, 2019) with 4 query heads sharing each KV head, SwiGLU (3 projections instead of 2), and RMSNorm. These represent the two major architectural families in production: NeoX-style (GPT-2/3, Pythia) and Llama-style (Mistral, Qwen, DeepSeek, Gemma). If the mechanism works on both, it is architecture-independent.

Run-003 used 40,000 steps (vs 20,000 for NeoX) because GQA with fewer attention parameters per layer converges slower. The Llama model has 7% fewer total parameters (405M vs 436M) due to the GQA reduction in KV projections. Both architectures reached comparable final PPL (~19-23 on overall training loss), confirming convergence despite the different training lengths.

5.2 Head Identification: Excess Scores

Identifying which heads specialize on delimiters requires correcting for the base rate of delimiter tokens in the probing text. A naive threshold (“heads where >50% of attention goes to delimiters”) is biased by format: JSON has 75.7% delimiter positions, so a uniform-attention head that distributes attention randomly across all positions would score 0.757. Without correction, this inflates head counts (168 apparent delimiter heads on NeoX, vs 50-66 after correction).

We use excess delimiter attention: the raw attention fraction minus the base rate. A head with 0.60 raw attention on a text where 0.45 of positions are delimiters has an excess score of 0.15: it attends to delimiters 15 percentage points more than chance. A head with 0.80 raw on JSON (0.757 base rate) has excess 0.043: it barely exceeds chance despite its high raw score.

Each head is probed on 4 texts spanning different delimiter densities: GCF generic (35% delimiter positions), GCF graph (38%), JSON (76%), and YAML (42%). The excess scores are averaged across all 4 texts. This averaging reduces sensitivity to any single format’s delimiter density and identifies heads that consistently specialize across formats.

Primary threshold: 0.15 excess. We also report counts at 0.10 and 0.20 for sensitivity analysis. On Llama, counts range from 85 (threshold 0.10) to 31 (threshold 0.20), with the primary 0.15 yielding 66. The causal findings (necessity, sufficiency, cross-format transfer) are internally consistent across all three thresholds.

5.3 Ablation Method

Our ablation methodology follows the zero-ablation approach established in the mechanistic interpretability literature (Olsson et al., 2022; Michel et al., 2019): for each ablation, we deep copy the model, zero the output projection weights for selected heads, measure per-format perplexity on held-out test data, then discard the copy. The original model is never modified. Each ablation is an independent measurement.

Why zero-ablation, not mean-ablation or resampling. Zero-ablation (setting output weights to zero) completely removes a head’s contribution to the residual stream. Mean-ablation (replacing ac-

tivations with their mean) and resampling-based methods (Conmy et al., 2023) preserve the head’s average contribution, which can mask causal effects when the head’s value comes from its variance across positions rather than its mean. For delimiter heads, the signal is precisely positional: they attend differently to delimiter vs content positions. Mean-ablation would preserve a diluted version of this signal, understating the causal effect. Zero-ablation is the stronger test.

Controls. Every delimiter head ablation is paired with a random head control: the same number of randomly selected non-delimiter heads are ablated, and the same metrics are measured. The causal signal is the gap between delimiter and random ablation, not the absolute direction of either. This controls for the generic effect of reducing model capacity. In some cases (particularly on Llama with GQA), removing any set of heads produces a regularization effect that improves performance on some formats. The delimiter-vs-random gap isolates the specific contribution of delimiter specialization from this generic capacity effect.

Architecture-specific considerations. On NeoX (full multi-head attention), each head has its own Q, K, V, and output projections. Zeroing the output projection fully disables the head; no residual signal leaks through. On Llama with GQA (4:1 ratio), each KV head is shared by 4 query heads. Zeroing one query head’s output leaves 3 siblings still using the same KV projection. The KV head still computes keys and values; the ablation only removes one of four interpretations of that shared signal. This makes per-query-head ablation a fundamentally weaker intervention on Llama than per-head ablation on NeoX. To address this, we developed KV-group ablation: zeroing all 4 query heads sharing one KV head simultaneously, which fully removes the KV head’s contribution and matches the intervention strength of NeoX’s per-head ablation (Section 8.3).

Ablation phases. The full ablation protocol consists of 18 phases on NeoX (run-002) and 12 phases on Llama (run-003), covering necessity, sufficiency, layer-wise, progressive scaling, individual head ranking, cross-format transfer at three thresholds, attention pattern extraction, per-token loss decomposition, embedding space analysis, bootstrap confidence, and controls (random ablation, transplant, KV-group).

5.4 Corpus

Both tokenizers were trained on the same corpus, and both models within each pair were pre-tokenized from the same source data. The corpus was designed to include structured data, code, and natural language in proportions that ensure the model sees all three categories without overweighting any:

Source	Size	%	Purpose
FineWeb (web text)	2.0 GB	33%	General language, diverse topics
Code (Go, Python, TS, JS, Rust)	800 MB	13%	Code syntax with delimiter characters

Source	Size	%	Purpose
JSON	850 MB	14%	Structured data in the dominant format
GCF	500 MB	8%	Structured data with clean delimiters
YAML/CSV/TOML	200 MB	3%	Additional structured formats
Natural language (Wikipedia)	200 MB	3%	Formal prose baseline

Total: ~4.5 GB. The corpus is deliberately not JSON-heavy: at 14%, JSON is a minority of the training data. This ensures that any advantage on structured data comes from the tokenizer’s treatment of delimiters, not from disproportionate training exposure.

Pre-tokenized token counts differ slightly between tokenizers (1,258,728,671 for structok-64k vs 1,269,271,190 for standard-64k, a 0.8% difference) because merge barriers produce more tokens on structured data and fewer on some natural language patterns. Both models see the same source text; only the tokenization differs.

5.5 Held-out Test Data

Product records with 6 fields (orderId, customer, status, total, date, category) at 5 sizes (5, 10, 20, 50, 100 records), generated with a different random seed (99999) from the training data generator. The held-out data uses different field values, different orderings, and different distributions than any training sample. Both JSON and GCF encodings of identical underlying data are generated, ensuring that format-comparison measurements reflect tokenization differences, not data differences.

Additional test sets include: GCF graph payloads (10 and 20 symbols with edges), alternative schemas (users, logs, API responses), code samples (Python, Go, TypeScript), adversarial payloads (pipe-like characters in values, embedded JSON syntax, numeric-heavy fields, empty fields), and 9 unseen formats for cross-format transfer evaluation (CSV, INI, SQL, Markdown tables, S-expressions, Protobuf text, TOML, TOON, XML).

6. Results: Baselines and Whole-Model Improvement

6.1 Baselines (Layer 1: Tokenizer)

Format	NeoX A	NeoX B	Ratio	Llama A	Llama B	Ratio
GCF generic	9,719	447,664	46x	15,166	152,264	10x
JSON	5,784,279	25,188,322	4.4x	195,337	1,288,738	6.6x
YAML	11,328	58,950	5.2x	13,524	54,306	4.0x
Code	603	2,972	4.9x	341	2,652	7.8x
Natural language (NL)	2,027	1,375	0.7x (B wins)	2,088	1,538	0.7x (B wins)

The merge-barrier model wins on structured data and code on both architectures by 3-46x. Natural language shows a small disadvantage (0.7x on this test set), though a separate Wikipedia evaluation (Section 6.3) shows NL perplexity tied at 19.4 vs 19.5. The structured data advantage is smaller on Llama (10x vs 46x for GCF), traced to GQA moderating the capacity for delimiter specialization (Section 8).

External validation on production models The comprehension gap introduced in Section 3.7 comes from an independent evaluation on production models (not the structok-trained models): 25 adversarial runs on 10 frontier models (Claude, GPT-5.5/5.4, Gemini), using 13 questions on 500-symbol payloads with 200 edges, zero format instructions. GCF **91.2%**, TOON 68.8%, JSON **54.1%**. On standard workloads (500 nested orders, 6 frontier models): GCF achieves 100% accuracy. These production-model results are consistent with the mechanistic findings from our controlled experiment: the tokenizer determines whether structural boundaries are visible, and production models with corrupted boundaries fail at scale.




Figure 8: Figure 8: GCF PPL scaling curve (fine-grained 8-size evaluation)

NeoX core evaluation (run-002) On held-out test data, the merge-barrier model (Model A) achieves 3x lower GCF perplexity:

Records	Model A GCF PPL	Model B GCF PPL	Advantage
5	1,900	3,642	1.9x
10	2,717	4,767	1.8x
20	3,952	9,810	2.5x
50	5,856	21,183	3.6x
100	9,719	33,703	3.5x

Model A wins 5/5 sizes. Average GCF PPL: 4,829 vs 14,621 (3.0x). Average GCF next-token accuracy: 3.7% vs 2.5% (+48%).

The merge-barrier tokenizer also produces fewer tokens for the same data (14-19% fewer on GCF: 113 vs 131 at 5 records, 1,855 vs 2,288 at 100 records). This means the model sees more data per token, contributing to the comprehension advantage.

Model B reads JSON better (1.9x lower JSON PPL), as expected: standard BPE merges JSON delimiters into familiar tokens from training. But this advantage comes at the cost of structural understanding.




Figure 9: Figure 9: Training convergence: standard BPE converges faster, both reach same PPL

Training convergence Standard BPE converges approximately 25% faster per step on overall perplexity. Model B reached PPL ~21 at step 8,000; Model A reached the same at step 10,000. But both settled to identical final PPL by step 20,000 (19.4 vs 19.5). The slower convergence is consistent with the merge-barrier tokenizer producing more tokens per text: the model needs more steps to see the same effective amount of data.

Fine-grained scaling (NeoX) Tested at 1, 2, 3, 5, 10, 20, 50, 100 records:

Records	Model A GCF PPL	Model B GCF PPL	Ratio
1	2,358	9,619	4.1x
2	2,294	5,315	2.3x
3	1,613	3,318	2.1x
5	2,296	5,147	2.2x
10	1,932	6,616	3.4x
20	3,374	13,593	4.0x
50	5,883	26,887	4.6x
100	8,112	43,152	5.3x

Model A wins 8/8 sizes. (Note: PPL values differ slightly from the 5-size table above because this uses a different, independently generated test set with 8 sizes.) The advantage grows monotonically from 2.1x at 3 records to 5.3x at 100 records. Larger payloads contain more delimiter boundaries; more boundaries means more opportunities for standard BPE's fused tokens to confuse the model.




Figure 10: Figure 10: Code comprehension: 3-5x better with merge barriers

Code comprehension (NeoX) An unexpected finding: merge barriers improve code comprehension 3-5x. The barrier characters ({, }, (,), :, ;) that protect structured data delimiters also protect code syntax.

Language	Model A PPL	Model B PPL	Advantage
Python	543	2,686	4.9x
Go	1,404	4,183	3.0x
TypeScript	729	2,667	3.7x

Model A wins 3/3 languages. This was not an explicit design goal of merge barriers but falls out naturally because code uses the same delimiter characters as structured data formats.




Figure 11: Figure 11: All formats comparison: Model A wins 11/11

All formats tested (NeoX)

Category	Test	Model A PPL	Model B PPL	Advantage
Structured	GCF tabular (avg)	4,829	14,621	3.0x
	GCF graph (10 sym)	14,095	39,558	2.8x
	GCF graph (20 sym)	18,289	36,314	2.0x
	Users schema	13,607	695,922	51x
	Logs schema	14,422	722,297	50x
	API response	1,935	14,075	7.3x
Code	Python	543	2,686	4.9x
	Go	1,404	4,183	3.0x
	TypeScript	729	2,667	3.7x
Other formats	YAML	5,439	16,872	3.1x
	CSV	2,847	30,616	10.7x
Natural language	Wikipedia	1,029	1,033	1.0x (tied)
Unseen format	TOON (tab-separated)	18,091	41,188	2.3x

Model A wins 11/11 categories with structured/code advantage, ties on natural language.




Figure 12: Figure 12: Adversarial robustness: Model A wins 5/5

Adversarial inputs (NeoX) GCF payloads with deliberately ambiguous content values:

Test	Model A PPL	Model B PPL	Ratio
Normal GCF	893	13,649	15.3x
Pipe-like chars in values	1,086	8,053	7.4x
JSON syntax embedded in GCF values	395	9,610	24.3x
Numeric-heavy fields	678	9,549	14.1x
Empty/missing fields	352	6,598	18.7x

Model A wins 5/5. The JSON-like values test embeds `{"key": "value"}` as a GCF field value. Model A handles it (PPL 395) because merge barriers keep embedded JSON syntax from confusing the model. Model B cannot distinguish embedded JSON from actual structure (PPL 9,610).

Cross-format transfer (NeoX) TOON (tab-separated) was never in the training data. Tab is a barrier character.

Format	Model A PPL	Model B PPL
GCF	55,000	2,844,107
TOON	18,091	41,188
JSON	1,328,211	1,802,773

Model A is 2.3x better on a format it has never seen, because the tab merge barrier generalizes.

6.2 Architectural Reorganization




Figure 13: Delimiter head specialization: 105 vs 23 heads

As a first indicator of how merge barriers reshape the model’s internal organization, we counted attention heads where >50% of attention goes to delimiter tokens (raw threshold, before the excess-score correction applied in Section 7.1):

Metric	Model A (barriers)	Model B (standard)
Delimiter-majority heads	105 / 384 (27%)	23 / 384 (6%)
Top head delimiter attention	85.3%	79.4%
Avg delimiter attention score	0.362	0.235

Model A develops **4.6x more delimiter-majority heads** (heads where over half of attention goes to delimiter tokens). The model builds dedicated circuitry for parsing structure when delimiters are cleanly isolated tokens. This is not a surface-level effect; the transformer’s internal architecture reorganizes in response to merge barriers.

6.3 Per-Token Loss




Figure 14: Figure 14: Per-token loss: delimiters are 2.4x easier for Model A

We computed cross-entropy loss at every token position on a 10-order GCF payload:

Metric	Model A	Model B
Avg delimiter loss	6.10	14.81
Avg content loss	13.28	14.74
Delimiter/content ratio	0.46x (delimiters easier)	1.00x (equal difficulty)

Model A finds delimiters 2.4x easier to predict than content. Model B finds delimiters equally hard as content. Model B's top-5 highest-loss tokens are all pipe characters (|): the model literally cannot predict where structure goes.

This is the mechanistic explanation for the perplexity gap. Model A has learned that delimiters are predictable structural markers. Model B treats them as arbitrary content.

Per-token loss under ablation (both architectures)

Condition	Delimiter loss	Content loss	Ratio
NeoX A (baseline)	6.1	13.3	0.46x
NeoX A (ablated)	5.7	11.4	0.50x
NeoX B	14.8	14.7	1.00x

Condition	Delimiter loss	Content loss	Ratio
Llama A (baseline)	5.6	12.1	0.46x
Llama A (ablated)	5.4	11.4	0.48x
Llama B	11.9	13.4	0.88x

Ablating delimiter heads does NOT spike delimiter loss back to Model B levels on either architecture. The prediction advantage is a whole-model property, distributed across all parameters, not localized to the specialized heads. This is the critical null result that establishes the Layer 2/Layer 3 distinction in the causal hierarchy.

6.4 Embedding Space




Figure 15: Embedding space: delimiter tokens cluster 69% more cohesively

Metric	Model A	Model B
Delimiter tokens in vocab	22	1,463
Delimiter internal cosine similarity	0.166	0.098
Separation metric (internal - cross)	0.174	0.115

Model A has 22 delimiter tokens (each barrier character is its own token, never merged). Model B has 1,463 tokens containing delimiter characters (merged with content). Model A's delimiter embeddings

are 69% more cohesive, forming a distinct cluster in embedding space. The model has learned that delimiters are a coherent category, distinct from content tokens.

The separation metric (internal similarity minus cross-category similarity) is 51% higher for Model A (0.174 vs 0.115). This means Model A's delimiter embeddings are not just close to each other; they are also farther from content embeddings. The embedding geometry reflects the tokenizer's clean separation: when delimiters are always their own tokens, the model learns embeddings that encode "I am structure" as a shared property, rather than blending structural and semantic information into merged-token embeddings.

6.5 Grammar Attention at Scale (NeoX)

Orders	Model A grammar%	Model B grammar%
5	37.1%	24.9%
10	31.4%	23.4%
20	30.8%	21.2%
50	30.5%	20.5%
100	29.7%	18.1%




Figure 16: Figure 16: Grammar attention at scale: Model A maintains 50% more

Model A allocates 50% more attention to grammar tokens at every scale and resists grammar attention collapse:




Figure 17: Figure 17: Grammar attention collapse comparison

Model	Small scale (5-10)	Large scale (50-100)	Change
Model A	34.3%	30.1%	-4.2%
Model B	24.1%	19.3%	-4.8%

Both models show some decay, but Model A starts higher and stays higher. Compare to the Gemma 2B finding from Section 3.7 (30% to 8.6% collapse): merge barriers prevent the catastrophic collapse observed in pre-existing models.

Ablating delimiter heads increases entropy by +1.0% on NeoX and +5.7% on Llama. Embedding space cohesion ratio changes by -5.5% (NeoX) and +7.9% (Llama). Random ablation controls produce equivalent changes. Entropy, grammar attention, and embedding structure are properties of the whole model, not controlled by the specialized heads.

Token repetition at scale

Orders	Model A GCF repeat%	Model B GCF repeat%	Model A tokens	Model B tokens
5	35.9%	44.3%	64	79
10	54.6%	62.8%	119	145
20	67.0%	73.3%	227	285
50	78.0%	81.0%	567	704
100	83.9%	84.6%	1,167	1,423




Figure 18: Token repetition at scale

Model A has lower token repetition because merge barriers prevent delimiter characters from being absorbed into content tokens. Each | is always its own token ID, but field values have more variety because they are not fused with delimiters. Model A also produces 14-19% fewer tokens because isolated delimiters are single tokens rather than multi-byte merged tokens.

Per-layer entropy profile Entropy at each transformer layer on 20-order GCF input (selected layers):

Layer	Model A	Model B	Delta
0 (input)	6.61	6.89	-0.28
4	4.87	4.69	+0.18
8	4.95	5.25	-0.30
12	4.87	5.69	-0.82
16	2.35	3.17	-0.82
20	2.75	1.96	+0.79
23 (output)	4.18	3.05	+1.13

Model A achieves its lowest entropy at layer 16 (2.35); Model B at layer 20 (1.96). The models process structure at different depths. Model A focuses earlier, consistent with having explicit delimiter boundaries that require less processing to resolve.

6.6 Confidence Calibration

Metric	Model A	Model B
Delimiter confidence (avg softmax prob)	0.086	0.058
Content confidence (avg softmax prob)	0.031	0.029
Delimiter/content confidence ratio	2.77x	2.00x

Model A is 48% more confident when predicting delimiter tokens (0.086 vs 0.058 average softmax probability). Both models are more confident on delimiters than content (delimiters are more predictable: they recur in fixed patterns), but Model A's delimiter-to-content confidence ratio is 2.77x vs Model B's 2.00x. Model A has learned that delimiters are substantially more predictable than content; Model B treats them as only moderately more predictable.

This confidence gap is consistent with the per-token loss finding (Section 6.3): Model A's lower delimiter loss (6.10 vs 14.81) translates directly into higher softmax probabilities for the correct delimiter token. The model assigns more probability mass to the right answer because it has stronger internal representations of where delimiters belong.

6.7 Delimiter Prediction Accuracy

Test	Model A delimiter acc	Model B delimiter acc
GCF tabular	25.9%	23.4%
GCF graph	4.5%	0.0%
JSON	4.4%	8.2%

Model A predicts GCF delimiters more accurately. Model B predicts JSON delimiters better (expected: it sees merged delimiter tokens during training). On GCF graphs, Model B gets zero delimiter predictions correct.

Content accuracy is near zero for both models on all tests (0.0% for Model A, 0.0-1.1% for Model B), confirming that the difference between models is entirely in structural token prediction, not in content prediction. Both models struggle equally with predicting specific data values; the divergence is in whether they can predict where structure goes.

6.8 Generation Quality

Both models generated 15/15 valid continuations across 5 prompt types (GCF tabular, GCF graph, JSON, Python, Go). But the quality differs substantially.

Model A generates recognizable structure: - GCF tabular: pipe-separated fields with plausible values
- Go: `http.Error w r. ( ) ( , Method )` (syntactically plausible function calls)

Model B generates garbled fusions: - GCF tabular: .@. || . || _ | 58240RD35. | Anderson-Zara | .@. (delimiters collapsed with content) - Go: wWriteHeaderhttpErrorwWriteHeaderwrwrwrwrwrwrwrwr (repetitive collapse, no structural boundaries)

The difference in generation quality is consistent with the per-token loss finding: Model A has learned delimiter positions as predictable structural markers, so it generates them in plausible positions. Model B treats delimiters as arbitrary content, so they appear randomly.

Sections 6.2-6.8 establish that merge barriers improve the entire model's structured data processing. The measurable improvements (per-token loss, entropy, embeddings, grammar attention, generation quality) are distributed across all parameters, not localized to specialized heads. This sets the stage for Section 7, which isolates the specific contribution of the specialized heads from this whole-model improvement.

7. Causal Ablation: Specialized Heads (Layer 3)

7.1 Head Identification

Using the excess-score method (Section 5.2) with threshold 0.15, head counts differ substantially from the raw >50% counts in Figure 13:

Model	Heads (excess 0.15)	Fraction
NeoX A (structok)	50	13%
NeoX B (standard)	3 (non-functional)	0.8%
Llama A (structok)	66	17%
Llama B (standard)	35 (functional)	9.1%

On NeoX, merge barriers are required for specialization (50 vs 3). With the raw >50% threshold (prior to base-rate correction), the count is 105 vs 23, a 4.6x ratio (Section 6.2). Model B, trained on the same data with the same architecture, develops only 3 heads that cross the excess-score threshold, and ablation confirms they are non-functional: removing them actually improves GCF (-22.9%) and GCF graph (-52.0%), while JSON/YAML/Code/NL stay within +/-3%. The heads are noise, not structural specialists.

On Llama, GQA enables partial specialization even without barriers (35 functional heads on standard Llama). Ablating these heads on Model B0 (our notation for the standard-BPE Llama control, analogous to NeoX Model B) drops GCF PPL by 53.8%, proving they are causal. This is a genuinely surprising finding: GQA's shared KV projections provide implicit structural priors that enable delimiter specialization even when the tokenizer corrupts delimiter boundaries. On NeoX with separate KV projections, this capability does not develop. This is discussed further in Section 8.4.

7.2 Necessity

Ablating delimiter heads on NeoX:

Format	Delimiter delta	Control delta	Gap
GCF generic	+58.7%	-35.8%	+94.5pp
YAML	+16.7%	-57.5%	+74.2pp
JSON	-37.0%	-74.7%	+37.7pp
NL	+4.2%	+2.9%	+1.3pp

Key finding: Delimiter head ablation **hurts** GCF generic (+59%) and YAML (+17%), while random head ablation **helps** those same formats (-36%, -58%). The effect is in opposite directions. This is the core causal result: removing the specific heads that specialize on delimiters degrades structured data comprehension, while removing the same number of random heads actually improves it (a regularization effect from reducing model capacity). Natural language is unaffected by either ablation type (+4% vs +3%), confirming the heads are structural specialists, not general-purpose.

JSON improves when delimiter heads are removed (-37%): the format-adversarial mechanism (Section 7.5). The heads trust structural boundaries, which becomes harmful when those boundaries are corrupted. Removing the trust removes the harm.

Bootstrap confidence: +16.7pp delimiter-random gap on GCF, std 2.0%, all 5 seeds consistent. The signal is statistically robust across different randomly generated test data.

On Llama, per-query-head ablation produces weaker absolute effects due to GQA (Section 8.6), but the delimiter-vs-random gap remains. KV-group ablation (Section 8.3) shows the gap most clearly: YAML +34.0% (delimiter KV groups) vs -18.0% (random KV groups), a +52.0pp gap. Llama bootstrap shows the GQA weakening on GCF (-0.4pp gap, mixed direction) but clear signal on JSON (+69.0pp, std 9.8%, all seeds consistent).

7.3 Sufficiency




Figure 19: Figure 19: Sufficiency scaling

On NeoX, keeping only 50 delimiter heads (removing 334 others) produces **better** GCF PPL than the full 384-head model (-55.6% at 30 rows, -55.2% at 50 rows). This is a remarkable result: 13% of the model’s heads, working alone with the other 87% zeroed out, produce better structured data comprehension than the full model. Random sets of 50 heads produce +38% to -31% PPL. The delimiter heads are not just “useful”; they are better than the full model at structured data.

This holds at all tested scales: 30 rows (-93pp gap between delimiter-only and random-only), 50 rows (-90pp), 100 rows (-63pp), 200 rows (-41pp). The gap narrows with scale. This is informative, not a weakness. At larger scales, the other 314 heads start contributing more to structured data comprehension. The delimiter heads carry the bulk of the signal at small scales, but at large scales the whole model needs to participate. This is consistent with the causal hierarchy: the heads are the specialized expression of a holistic improvement, and at scale the holistic improvement matters more than the specialization. The gap never reverses: delimiter heads maintain a meaningful advantage even at the 2048-token context window limit.

On Llama, sufficiency is weaker due to GQA’s shared KV projections. Zeroing query heads in the reverse ablation doesn’t fully remove the delimiter signal because the shared KV projections still route information through surviving random query heads. The per-query-head intervention is too weak under GQA to cleanly demonstrate sufficiency, though delimiter-only still beats baseline at all sizes.

7.4 Layer-Wise Ablation




Figure 20: Figure 20: Layer-wise comparison NeoX vs Llama

Layer group	NeoX heads	NeoX GCF delta	Llama heads	Llama GCF delta
Early (0-7)	6	-10%	25	-41%
Middle (8-15)	14	+4%	25	+20.1%
Late (16-23)	20	+63%	16	+6.4%

On NeoX, late layers are causal (+63% GCF degradation). Early and middle layers barely matter. If delimiter specialization were simple pattern matching (“this token is a pipe character”), it would happen in early layers. The late-layer concentration indicates the model uses delimiters for abstract structural reasoning: parsing field boundaries, tracking record structure, navigating between sections. The model doesn’t just see delimiters; it reasons through them.

On Llama, the distribution is different: delimiter heads spread across early and middle layers (25/25/16), with middle layers showing the strongest causal effect (+20.1%). GQA pushes structural processing earlier because shared KV projections force structural information to be consolidated before it propagates through the query heads. This is an architectural effect, not a mechanism failure: the model still reasons through delimiters, it just does so at different depths.

7.5 Attention Saturation (Format-Adversarial Mechanism)



Figure 21: Figure 21: Attention heatmap GCF vs JSON

Head	Architecture	GCF c->content	JSON c->content
L17H1	NeoX	14.4%	0.9%
L6H0	Llama	35.3%	0.8%

JSON attention saturates identically on both architectures. Delimiter heads send 97-99% of content attention to delimiters on JSON, leaving under 1% for actual content. The model is structurally blind to the data it is supposed to reason about. On GCF, the same heads retain 14-35% content attention, enough to extract information. The difference: GCF has 35-38% delimiter positions while JSON has 76%. JSON's delimiters are everywhere, so attending to them is less informative. GCF's delimiters are sparse structural markers that carry high signal.

This is the format-adversarial mechanism: heads learn “attend to delimiters for structural reasoning” from training on GCF data, and when applied to JSON (where delimiters are inside merged tokens and delimiters are ubiquitous), that trust becomes actively harmful. The heads are not format-neutral; they are format-adversarial to corrupted formats. This is a novel mechanism: training-induced trust in structural boundaries becomes a liability when those boundaries are corrupted.

The structural pattern test (Appendix C) provides further confirmation. GCF's own pipe delimiter becomes actively adversarial (-54%) when placed in an unfamiliar wrapping layout. The heads learned “pipe means flat field separator.” When pipe appears in a wrapping context, they misapply that reasoning and actively hurt comprehension. Unfamiliar delimiters (tab) transfer in all layouts (+32% to

+123%) because no conflicting prior exists. The mechanism is about learned character-specific priors conflicting with context.

7.6 Adversarial Robustness Under Ablation

We tested whether delimiter heads contribute to detecting structural corruptions in GCF payloads. Five corruption types were tested: wrong delimiter (comma instead of pipe), missing field, extra pipe, wrong record count, and broken header. Detection is measured as the PPL spike ratio (corrupted/clean): values above 1.5x indicate the model notices the corruption.

On NeoX, ablating delimiter heads reduces structural corruption detection by ~56% across 3 of 4 corruption types:

Corruption type	Baseline spike	Ablated spike	Reduction
Missing fields	89.8%	29.8%	-67%
Wrong header	59.0%	15.6%	-74%
Swapped values	130.3%	40.2%	-69%
Wrong delimiter	61.1%	62.9%	+3% (retained)

Wrong-delimiter detection is fully retained under ablation (+62.9% vs +61.1%), suggesting that mechanism operates through a different pathway, likely the whole-model layer (Layer 2) rather than the specialized heads (Layer 3). The model detects “this character is not pipe” through its general delimiter embeddings, not through the attention pattern of specialized heads.

On Llama, the pattern is similar: heads contribute to but don’t solely control error detection, consistent with the causal hierarchy. Detection is partly a Layer 2 whole-model property (delimiter embeddings encode what delimiters look like) and partly Layer 3 (specialized heads notice when expected structural patterns are disrupted).

7.7 Head Ranking and Emergence

We ablated each delimiter head individually to measure per-head importance. On NeoX, using a broader threshold (0.10 excess) to capture edge cases, 39 of 74 candidate heads hurt GCF when removed (positive delta); 34 help (negative delta, threshold artifacts). The top 5 heads account for 45% of total degradation (44pp of 98pp). On Llama, 36 of 66 heads hurt GCF, top 5 account for 36%. Structural reasoning is concentrated in a small core of heads in late layers (NeoX) or middle layers (Llama), not distributed uniformly across all delimiter-specialized heads. The remaining heads that attend to delimiters but don’t contribute causally are likely threshold artifacts: they cross the >50% line but don’t carry structural reasoning.

Delimiter heads emerge early during training and narrow with continued training. On NeoX, probing checkpoints at threshold 0.10 to track the broader population: 107 heads at step 1,000 (from an earlier

training run with the same configuration), narrowing to 70 at step 3,500 and 61 by step 5,000, with concentration (fraction of total excess attention in the top 10% of heads) increasing from 50% to 54%. The model starts with many loosely specialized heads and prunes to a stable core. On Llama, probing from step 15K (same 0.10 threshold): 71 heads at step 20K, narrowing to 49 by step 35K. At the primary 0.15 threshold, the final counts are 50 (NeoX) and 66 (Llama); the narrowing trends track the same pruning dynamic at a broader threshold. There is no phase transition on either architecture. The model exploits clean delimiter tokens from the earliest stages of training; continued training prunes weak heads, increasing concentration. The convergent final counts (50-66 at 0.15 threshold) across both architectures despite different timelines suggest a convergent property of the mechanism.

8. Cross-Format Transfer and Architecture Independence

Section 6 established the whole-model improvement, and Section 7 proved that delimiter heads are necessary and sufficient for structured data comprehension on formats seen during training. The next question is whether these heads generalize: do they help on formats the model has never seen?

8.1 Universal Transfer (Layer 4)




Figure 22: Figure 22: Cross-format transfer NeoX vs Llama

Format	NeoX (50 heads)	Llama (66 heads)	Transfer?
CSV	+38.0%	+27.0%	YES
INI	+41.0%	+1.5%	YES/weak

Format	NeoX (50 heads)	Llama (66 heads)	Transfer?
SQL	+72.1%	+13.5%	YES
Markdown table	+20.9%	+38.0%	YES
S-expression	+25.3%	+23.8%	YES
Protobuf text	+120.4%	+17.9%	YES
TOML	+74.8%	+21.3%	YES
TOON	-2.7%	+15.3%	neutral/YES
XML	+12.2%	+34.0%	YES

8 of 9 unseen formats transfer on both architectures. Average degradation across all 9 formats when delimiter heads are removed: +44.7% on NeoX, +21.4% on Llama. Excluding the one non-transferring format on each architecture (TOON on NeoX, INI on Llama): +50.6% and +23.9%. The transfer spans every delimiter style: commas (CSV), equals signs (INI), parentheses (SQL, S-expressions), pipes (Markdown tables), braces and colons (Protobuf). The mechanism is not specific to pipe or to GCF; it is a general property of delimiter specialization.

Seven hypotheses were tested over the course of the study to explain apparent selectivity in early results, which showed 6 of 9 formats transferring: delimiter density ($r=0.026$, $p=0.927$), merge word count, merge rate, structural pattern, boundary clarity ($r=-0.29$, $p=0.44$), positional distribution ($r=+0.08$, $p=0.83$), and spacing regularity ($r=-0.14$, $p=0.71$). None were predictive. The resolution came from a methodological discovery: the same tab-separated format showed -15.8% with 76 heads (one identification run) and +32.1% with 88 heads (a different identification run), with nearly identical baselines (21,471 vs 21,223 PPL). The format didn’t change; the head set did. Transfer selectivity was an artifact of head identification instability, not a property of the formats. With corrected excess-score identification, transfer is universal.

8.2 What Replicates Across Architectures

Seven findings replicate across both architectures despite the substantial differences between NeoX (learned positions, full MHA, GELU, LayerNorm) and Llama (RoPE, GQA 4:1, SwiGLU, RMSNorm):

Finding	NeoX	Llama	Status
Delimiter heads emerge	50 heads (13%)	66 heads (17%)	Replicates
Cross-format transfer	8/9 formats	8/9 formats	Replicates
JSON attention saturation	99.1% (L17H1)	99.2% (L6H0)	Replicates (nearly identical)
Head ranking concentration	Top 5 = 45%	Top 5 = 36%	Replicates

Finding	NeoX	Llama	Status
Head count narrowing (0.10 threshold)	107->61 (steps 1K-5K)	71->49 (steps 20K-35K)	Replicates
Per-token loss is whole-model	Ablation: 0.46x->0.50x	Ablation: 0.46x->0.48x	Replicates
Natural language unaffected	+4.2% delimiter, +2.9% random	Similar	Replicates

The replication is not trivial. GQA changes the magnitude of every effect (10x vs 46x baselines, +21.4% vs +44.7% average transfer), the layer distribution (middle vs late), and the B model’s baseline capability (35 functional heads vs 3 non-functional). But the qualitative findings are identical: merge barriers cause delimiter head emergence, those heads are necessary for structured data comprehension, they transfer to unseen formats, and the underlying prediction advantage is a whole-model property. The mechanism is architecture-independent; only its expression varies with attention design.

8.3 KV-Group Ablation (New Methodology for GQA)

The trained-format ablation results on Llama showed direction flips compared to NeoX (GCF improved under ablation instead of degrading). The direction flips raised the question: is the mechanism failing on Llama, or is the ablation intervention too weak under GQA?

The answer is the latter. Per-query-head ablation under GQA is fundamentally different from per-head ablation under full MHA. On NeoX, zeroing a head’s output projection removes its entire Q/K/V/O contribution. On Llama, zeroing one query head leaves 3 siblings still using the same KV projection. The KV head still computes the same keys and values; the ablation only removes one of four interpretations of that shared signal. The delimiter information persists through the surviving siblings.

We developed KV-group ablation as the GQA-equivalent of NeoX’s per-head ablation: zero all 4 query heads sharing one KV head simultaneously. This removes the KV head’s entire contribution, matching the intervention strength of NeoX’s per-head ablation. 19 delimiter KV groups were identified out of 96 total (20%).




Figure 23: Figure 23: KV-group ablation gaps

Format	Delimiter KV delta	Random KV delta	Gap
GCF	-15.5%	-63.9%	+48.4pp
JSON	+67.8%	-96.2%	+164.0pp
YAML	+34.0%	-18.0%	+52.0pp

The delimiter-vs-random gap is the causal signal. Removing delimiter KV groups is dramatically worse for structured formats than removing random KV groups. The +164pp JSON gap means delimiter KV groups are specifically important for JSON processing: they are format-adversarial (Section 7.5), and removing them from the delimiter set frees JSON from the adversarial effect less than removing random groups does. GCF's negative absolute delta (-15.5%) is a regularization effect from removing 20% of KV groups, but the +48pp gap proves delimiter groups are specifically important compared to random groups.

Under GQA, the correct metric for causal ablation is the delimiter-vs-random gap, not the absolute direction. This is a methodological contribution applicable to any GQA-based ablation study.

8.4 The B0 Finding

The most surprising result from run-003: standard Llama (Model B0, trained with standard BPE, no merge barriers) develops 35 functional delimiter-specialized heads. This is fundamentally different from NeoX, where Model B developed only 3 non-functional heads. As noted in Section 7.1, ablating B0's 35 heads confirms they are causal (53.8% GCF PPL drop), not artifacts of the identification method.

The B0 finding reframes the story from “merge barriers create something that doesn’t otherwise exist” to “merge barriers amplify a capability that GQA partially enables.” GQA’s shared KV projections act as implicit structural priors: because 4 query heads must share the same key-value representation, the model is forced to develop KV heads that capture general structural patterns (including delimiter boundaries) rather than allowing each head to independently ignore them. On NeoX with separate KV projections, each head can independently learn to ignore delimiters, and without merge barriers, they do.

The implication for production models is significant. Llama-family models (Mistral, Qwen, DeepSeek, Gemma) all use GQA. They likely already have partial delimiter specialization, which may explain their reasonable structured data performance despite corrupted tokenization. Merge barriers would amplify this existing capability (66 heads vs 35, a 1.9x ratio). On NeoX-family models, the improvement would be larger (50 vs 3, a 16.7x ratio) because the baseline is near zero.

8.5 Stranded Heads: A Previously Unknown Attention State

Every production model that processes structured data pays a cost nobody has measured: attention capacity allocated to boundary recovery that produces no useful output. We identify this cost as a new attention state, distinct from both productive specialization and dormancy, and demonstrate it is caused entirely by the tokenizer and reversible without retraining.

The B0 finding raises a question: if standard NeoX develops only 3 non-functional delimiter heads, are these heads incapable of structural anchoring, or merely stranded by corrupted input? We test this by feeding Model B (standard NeoX) the SAME data in two formats: JSON (boundaries merged into tokens) and GCF (boundaries preserved as clean pipe tokens). Same weights, only input format differs.

Result: 16 stranded heads identified. These heads show significantly focused attention on GCF’s clean delimiters (low entropy, high delimiter mass) but diffuse attention on JSON’s merged tokens (high entropy, low delimiter mass). The model is CAPABLE of structural anchoring but STRANDED by the tokenizer.

Metric	JSON input	GCF input
Mean attention entropy (384 heads)	6.38	5.92
Heads more focused (lower entropy)	—	344/384 (89.6%)
Strongly stranded (drop > 0.3, mass > 0.03)	—	16 heads
Strongest head (L1H7) delimiter mass	46.1%	82.5%

Position correspondence with Model A. Two of the 16 stranded heads occupy the exact same layer/head position as Model A’s top structural anchoring heads:

- **Layer 1, Head 5:** Model A anchoring score 0.58 / Model B stranded (entropy drop 0.79)
- **Layer 18, Head 11:** Model A anchoring score 0.57 / Model B stranded (entropy drop 0.39)

The same architectural slot succeeds (Model A, clean boundaries) or strands (Model B, merged boundaries). The only difference is the tokenizer.

89.6% of all heads are more focused on GCF (not just the 16 stranded ones). This confirms the whole-model boundary recovery cost described in Layer 2 of the causal hierarchy: the standard tokenizer degrades attention focus across the entire model, not just in specialized heads.

Interpretation: a third state. Prior work describes two head states: functional (productive, necessary) and dormant (collapsed into attention sinks, safely removable; Sandoval-Segura et al., 2025). We identify a third state: **stranded**. Stranded heads are active but unproductive, spreading attention diffusely across the sequence (high entropy) rather than concentrating on P0 (dormant) or on delimiters (anchoring). They represent the worst possible outcome for compute efficiency: they cannot be safely zeroed (the model is still attempting structural processing through them), yet they produce no useful structural signal.

The three states form a spectrum determined by tokenizer-attention coupling:

State	Attention pattern	Productive?	Safely removable?	Cause
Structural anchoring	Concentrated on delimiters	Yes	No	Clean boundaries (merge barriers)
Stranded	Diffuse, high entropy	No	No	Corrupted boundaries (standard BPE)
Dormant	Concentrated on P0 sink	No	Yes	No structural work needed

To our knowledge, this intermediate state has not been previously identified. The prior taxonomy (Sandoval-Segura et al., 2025; Gu et al., 2025) treats heads as either functional or dormant. Stranding is a third state that is computationally worse than either: dormant heads at least route to a near-zero value vector, producing minimal disruption and safely removable. Stranded heads distribute attention across hundreds of tokens, each contributing noise to the output, while remaining unsafe to remove because the model is still attempting structural processing through them.

The stranding tax. The 89.6% figure (344/384 heads more focused on GCF than JSON) implies that stranding is not limited to 16 strongly-stranded heads. The entire model operates suboptimally on structured input with corrupted boundaries. Every API call that includes JSON is paying a stranding tax: attention capacity spread diffusely instead of concentrated productively. This cost scales with

the proportion of structured data in the input and the number of delimiter positions that BPE has corrupted.

Reversibility without retraining. The position correspondence finding (same architectural slot succeeds in Model A, strands in Model B) combined with the format-switching result (same weights, different input, different behavior) demonstrates that stranding is reversible at inference time. A model does not need to be retrained with merge barriers to benefit from clean boundaries; it needs only to receive input where structural delimiters are already preserved as distinct tokens. This is exactly what GCF provides as a wire format: clean delimiter tokens that unstrand heads without model modification.

Training dynamics confirm frustration. We trained fresh 410M model pairs (2000 steps, gradient measurements every 50 steps) and measured the correlation between gradient signal pushing toward delimiter attention and actual delimiter attention achieved. Model A (merge barriers) shows strong gradient-attention coupling ($R=-0.58$, $p=0.0001$): as heads learn to attend to delimiters, gradient magnitude decreases (the learning is succeeding, less correction needed). Model B (standard BPE) shows weaker coupling ($R=0.39$, $p=0.012$): the gradient keeps pushing but attention responds slowly. The tokenizer determines whether gradient signal translates to specialization. This is stranding caught in the act during training: the loss function wants delimiter attention, the model tries, and the corrupted input defeats the attempt.

Implication for the dormant heads literature. Sandoval-Segura et al. (2025) asked whether “a single measure can accurately predict the fraction of dormant heads.” Our finding suggests the question itself is incomplete: the dormant/active binary misses the stranded middle. A complete characterization of attention head states requires three categories, and the transition between them is governed by the tokenizer’s treatment of structural delimiters in the input.




Figure 24: Figure 24: Gradient-attention coupling: same training objective, different outcomes. Model A ($R=-0.58$) responds to gradient signal; Model B ($R=0.39$) stalls.




Figure 25: Figure 25: Stranded heads entropy scatter. Red = 16 stranded heads (entropy drop > 0.3). Cyan = Model A anchoring positions.




Figure 26: Figure 26: Position correspondence. Green circles mark exact matches: same slot anchors in Model A, strands in Model B.




Figure 27: Figure 27: 344/384 heads (89.6%) are more focused on GCF than JSON with identical weights.

8.5.1 Definitive Proof: Tokenizer Swap on Frozen Weights The format-comparison probe above has a confound: JSON and GCF are different formats with different token sequences, so attention differences could reflect format properties rather than delimiter cleanliness. We eliminate this confound with forced-clean tokenization.

Method. Segment the input text at the 16 barrier characters, then tokenize each segment independently with standard-64k’s BPE encoder. Every resulting token ID is native to Model B’s vocabulary. No remapping, no vocabulary mismatch. The only difference between the “normal” and “clean” conditions is whether "name is one token or two (" + name).

Result: 384/384 heads wake up. Every attention head in Model B shows dramatically more delimiter attention when given clean token boundaries. Average delimiter attention mass rises from 14% (normal) to 55% (clean), a delta of +0.40. The minimum delta across all 384 heads is +0.35. No head decreases its delimiter attention under clean tokenization.

The 16 previously identified stranded heads wake up with average delta +0.39, indistinguishable from the non-stranded average of +0.40. The effect is universal: every head in the model has latent structural capacity that clean tokenization unlocks.




Figure 28: Figure 28: Tokenizer swap on frozen weights. Each vertical line connects one head’s delimiter attention under normal tokenization (blue, 14%) to clean tokenization (green, 55%). Red lines highlight the 16 previously identified stranded heads. All 384 lines point upward.

This result changes the interpretation of stranding. Stranded heads are not a small subpopulation. Every head is stranded. The original 16 were the most visibly affected under the less sensitive format-comparison probe. The true stranding population is the entire model.

8.5.2 Stranding Is Immediate and Permanent Method. Run the tokenizer-swap analysis at each of 8 intermediate checkpoints saved during Model B’s training (Llama 410M, standard BPE, steps 5,000 through 40,000 at 5,000-step intervals).

Result: both lines are flat.

Step	Normal mass	Clean mass	Delta	Heads woke up
5,000	14.3%	54.1%	+39.8 pp	384/384
10,000	14.4%	54.2%	+39.9 pp	384/384
20,000	14.3%	54.2%	+39.9 pp	384/384
40,000	14.4%	54.2%	+39.9 pp	384/384

The frustration gap (41 percentage points of unrealized delimiter attention) is present from step 5,000 and does not change across 35,000 additional steps of training. The model builds full structural capacity early and the tokenizer locks it out permanently.




Figure 29: Figure 29: The frustration gap across training. Green line: what the model CAN do (55% delimiter attention under clean tokenization). Blue line: what the tokenizer ALLOWS (14% under normal tokenization). The 41pp gap is flat from step 5,000 through step 40,000.

Training cannot overcome stranding. The gradient-attention coupling data ($R=-0.58$ for Model A vs $R=0.39$ for Model B) confirms that gradient signal toward delimiter specialization exists but cannot translate into attention pattern changes when token boundaries are corrupted. The gradient pushes; the tokenizer blocks; the gap persists.

8.5.3 P0 Drift: Stranding Is Stable, Not a Pathway to Dormancy A natural hypothesis connects stranding to dormancy: stranded heads that cannot find productive anchoring targets might progressively collapse into position-0 attention sinks (Sandoval-Segura et al., 2025; Gu et al., 2025). If true, stranding would be transient and the wasted capacity eventually recoverable through head pruning.

Result: no drift. P0 mass under normal tokenization is flat at 14.0% from step 5,000 through step 30,000. P0 mass under clean tokenization is flat at 11.4%. The gap (+2.5pp) is stable at every check-point. Stranded heads lean slightly toward the P0 sink under corrupted input but do not collapse.




Figure 30: Figure 30: P0 drift analysis. No drift toward dormancy across training. Stranded heads show a stable +2.5pp P0 bias but never fall into the sink.

The three-state spectrum (anchoring, stranded, dormant) describes genuinely distinct states, not a continuum. Stranded heads do not become dormant with more training. The wasted capacity is permanently locked: neither productive nor removable.

8.5.4 Stranding at 1.3B Scale The tokenizer-swap experiment replicates at 1.3B: **768/768 heads wake up** (delimiter mass 14% to 55%, delta +0.40, minimum delta +0.35). The universality of stranding scales with the model.

Model B at 1.3B develops 124 delimiter heads (16.1% of 768). These heads are not merely non-functional; they are counterproductive. Removing them improves GCF comprehension by 57%. The damage is consistent across all scales:

Scale	Arch	Model B heads	GCF ablation delta
410M NeoX	MHA	3	-22.9% (removal helps)
410M Llama	GQA	35	-53.8% (removal helps)
1.3B Llama	GQA	124 (16%)	-57.0% (removal helps)




Figure 31: Figure 31: The scaling paradox of standard BPE. Left: both models develop more delimiter heads at scale, but only Model A’s heads are functional. Right: Model B’s heads are counterproductive at every scale, and the damage grows.

More capacity invested in delimiter detection on corrupted boundaries means more capacity wasted on counterproductive attention patterns. Scaling does not solve stranding. It amplifies it.

8.6 GQA Effects (Not Mechanism Failure)

Every difference between NeoX and Llama results traces to GQA’s shared KV projections:

Effect	Cause
Smaller A/B baseline ratio (10x vs 46x)	Fewer attention parameters per layer, less capacity to exploit clean delimiters
Earlier layer distribution (middle vs late)	Shared KV forces structural consolidation before information propagates to query heads
Weaker per-head ablation effects	3 siblings still use the same KV signal after one query head is zeroed
B0 has functional heads (35 vs 3)	Shared KV provides implicit structural priors that separate-KV cannot
Lower concentration (13-15% vs 37-54%)	Specialization spreads across more query heads per KV group
Trained-format direction flips	Regularization from removing many heads competes with weaker causal signal

None of these indicate the mechanism fails. They indicate the ablation methodology needs adaptation for GQA (use KV-group ablation and gap-based metrics), and that GQA changes how the mecha-

nism expresses itself, not whether it exists. The core finding is architecture-independent: clean delimiters cause concentrated head specialization on both architectures.

9. The Causal Hierarchy

The experiments reveal a four-layer causal hierarchy. Every attempt to localize a metric to the delimiter heads showed it was a whole-model property instead. The per-token loss ablation (null: heads don't control the 2.4x advantage), the attention entropy ablation (null: heads don't control entropy patterns), the embedding space ablation (null: heads don't control embedding cohesion), and the transplant controls (any Model A weights help Model B) all point to the same conclusion. The hierarchy, from root cause to observable effects:

Layer 1: Tokenizer (root cause). Clean delimiters vs corrupted. This is the only variable in the controlled experiment. Everything else flows from it. The tokenizer shapes every weight during training: embeddings, attention projections, feed-forward layers, layer norms. The merge barrier is a pre-tokenization constraint; its effects propagate through the entire training process.

Layer 2: Whole model (first-order effect). Better embeddings, better per-token prediction (2.1-2.4x delimiter advantage), lower attention entropy, 3-46x overall PPL improvement on structured data. These are properties of the entire model, not localized to any specific heads. Ablating 50-66 heads doesn't change these properties because they're distributed across all 405-436 million parameters. The tokenizer improvement is holistic: clean delimiters give the model a cleaner gradient signal at every training step, which shapes all weights, not just the ones that end up in delimiter-specialized heads.

Layer 3: Specialized heads (second-order effect). 50-66 delimiter-majority heads emerge immediately (by step 1,000 on NeoX) and sharpen with training. They are causally necessary for format-level comprehension (ablating them hurts GCF +59%). They are sufficient: 13% of heads alone beat the full model on structured data. They concentrate in late layers on NeoX (reasoning, not pattern matching) and middle layers on Llama (GQA effect). They become format-adversarial to corrupted formats (JSON improves when they're removed because they trust corrupted boundaries). But they don't control per-token loss, entropy, or embeddings. They are the specialized expression of the whole-model improvement, not its source.

Layer 4: Cross-format transfer (third-order effect). The specialized heads generalize to 8 of 9 unseen formats (+44.7% average degradation on NeoX, +21.4% on Llama). The transfer spans every delimiter style tested (commas, parentheses, braces, pipes, equals signs). Seven hypotheses were tested to explain apparent selectivity; the resolution was that selectivity itself was an artifact of head identification instability. Transfer is universal, not selective. The mechanism learned from GCF's pipe delimiters applies to any format with clean structural boundaries.

The distinction between layers is critical for interpreting the results correctly. Leading with the heads and treating them as "the mechanism" overstates their role (layer 3). The mechanism is the tokenizer

(layer 1). The heads are the most visible and ablatable evidence of the mechanism, but they are not its source.

Transplant experiments (Appendix D) make this distinction concrete: transplanting Model A’s delimiter heads into Model B improves GCF by 81%, but transplanting the same number of random non-delimiter heads also improves it by 70%. On JSON and TOON, random heads improve Model B *more* than delimiter heads. The improvement is holistic: merge barriers improve all of Model A’s weights through training, and the transplant reflects this general quality improvement rather than head-specific portability. The ablation (within-model, causal) remains the stronger evidence because it is unconfounded by the overall quality difference between models.

10. Scaling to 1.3 Billion Parameters

The experiments in Sections 5 through 9 established the merge barrier mechanism at 410M parameters on two architectures. Run-004 tests whether delimiter head specialization is a small-model artifact that vanishes at scale, or whether it strengthens as model capacity grows.

10.1 Training Configuration

The 1.3B Llama uses 24 layers with 32 attention heads per layer (768 total heads), compared to the 410M models’ 24 layers with 16 heads per layer (384 total). All other hyperparameters match the 410M Llama configuration, with two exceptions necessitated by numerical stability at scale: precision was switched from fp16 to bf16 after NaN gradients appeared at step 11,710, and the learning rate was halved to 1.5e-4 with gradient clipping set to 0.5. The training corpus, tokenizer, and evaluation protocol are identical to the 410M experiments. Both models trained for 50,000 steps on a single A100 80GB GPU.

10.2 A/B Comparison

Format	Model A (barriers)	Model B (standard)	Ratio (B/A)
GCF generic	3,368	2,484,060	738x
JSON	104,241	289,437	2.8x
Python	277	N/A	N/A

The headline number at 1.3B is **738x** on GCF, dramatically larger than the 46x observed at 410M NeoX. Model B essentially cannot parse GCF (PPL 2.5 million). Even on JSON (which both models saw during training), Model A is 2.8x better.

10.3 Delimiter Head Identification

Model A develops 161 delimiter heads out of 768 total (21.0%). Cross-format transfer testing confirms 8/9 unseen formats experience degraded perplexity when these 161 heads are ablated, with an average degradation of +420.5%.

Model B develops 124 delimiter heads (16.1%). Their ablation effect is discussed in Section 8.5.4.

10.4 Scaling Trend

Metric	410M GPT-NeoX	410M Llama	1.3B Llama
Delimiter heads (Model A)	50	66	161
Specialization fraction	13.0%	17.2%	21.0%
Cross-format transfer	8/9	8/9	8/9
Avg degradation (Model A)	+59%	+21.4%*	+420.5%
Delimiter heads (Model B)	3	35	124
Model B ablation effect	-22.9%	-53.8%	-57.0%

The 410M Llama degradation is lower than NeoX because GQA enables partial specialization without barriers (the B0 finding). Within the Llama architecture, degradation increases 20x from 410M to 1.3B.

Two trends emerge. First, the specialization fraction increases: 13% to 17% to 21%. Larger models allocate a greater share of attention capacity to delimiter processing. Second, within the Llama architecture, the causal importance of Model A’s heads increases dramatically (20x), while Model B’s heads become increasingly counterproductive (-53.8% to -57.0%). Scaling does not solve the stranding problem. It amplifies it.

11. Domain Generalization

The preceding sections established the merge barrier mechanism on structured data formats. This section tests whether the same mechanism operates in fundamentally different domains.

11.1 Code Domain

We trained a 410M GPT-NeoX model pair for 20,000 steps on a 4GB subset of The Stack v1 deduplicated (Python 35%, Go 15%, TypeScript 15%, Rust 15%, Java 20%). The barrier character set was adapted for code: () . < > { } , = : @ ; (12 characters, selected by measuring merge rates above 50% on o200k_base).

Model A achieved overall code perplexity of 4.9 compared to Model B's 7.3, a **1.5x improvement**. 83 delimiter heads were identified (22% of 384), with 7/9 unseen formats hurt by ablation.

The strongest evidence comes from per-token loss analysis:

Delimiter	Merge Rate	B/A Ratio
(paren	100%	1.5x
. dot	90%	1.2x
, comma	70%	1.2x
{ brace	67%	0.9x (reversal)
= equals	60%	0.7x (reversal)

Characters with higher merge rates show larger improvements. Characters below ~65% merge rate show reversals, where the compression benefit of BPE outweighs the corruption cost. This monotonic relationship between merge rate and benefit is predicted by the merge barrier hypothesis and is not predicted by any alternative explanation. It cannot be caused by vocabulary size, training dynamics, or any global property, because those would affect all token types uniformly.

Function calls (2.1x) and member access (1.7x) show the largest category-level improvements, consistent with their dependence on the most heavily merged characters (parentheses, dots). The code delimiter head fraction (22%) exceeds the structured data fraction (13%) despite using fewer barrier characters (12 vs. 16), reflecting that code requires more diverse structural parsing patterns (function calls, member access, type annotations, decorators, scope boundaries).

11.2 Chemistry Domain

We trained a 410M GPT-NeoX model pair for 20,000 steps on ZINC-250K (249,455 drug-like molecules with 10x SMILES augmentation, 113 MB). The barrier character set: () = / [] (6 characters).

Model A achieved SMILES perplexity of 18.1 vs Model B's 39.2, a **2.2x improvement**. Model A generated **67% more valid molecules** (40/500 vs 24/500) and committed **2.5x fewer valence errors** (4.2% vs 10.4%). 87 delimiter heads were identified in Model A vs 69 in Model B.

The absolute validity rates (8% and 4.8%) are low, expected for 410M models trained for 20,000 steps on a 113 MB corpus, not production molecular generators. The finding is the relative improvement, not the absolute rate.

Extending Jindal and Ju (2026). Jindal and Ju's SMolLM used a character-level tokenizer (vocabulary size 30) to train a 53K-parameter SMILES model. A character-level tokenizer is mechanically identical to merge barriers: every character is its own token, so brackets never merge with adjacent atoms. They found a single bracket-matching attention head at that tiny scale without understanding why. Our experiment explains their result: their tokenizer was an implicit merge barrier, and their bracket-matching head was structural anchoring enabled by clean boundaries. We find 87 delimiter heads at 410M parameters, following the same scaling pattern.

11.3 Cross-Domain Analysis and Tokenizer-Attention Coupling

Three domains, three corpora, three barrier character sets, one consistent result:

Domain	PPL Ratio (A/B)	Delimiter Heads (A)	Key Finding
Structured data	46-738x	50-161 (13-21%)	Causal ablation proof, scales with model size
Code	1.5x	83 (22%)	Per-token advantage maps to merge rate
Chemistry	2.2x	87 (23%)	67% more valid molecules, extends Jindal

Effect magnitude correlates with delimiter density. The fraction of attention heads that specialize on delimiters is stable at 13-23% across all experiments. The direction is always consistent: merge barriers improve perplexity on delimiter-dependent structures without degrading natural language or non-delimiter content.

Layer distribution varies by domain. On NeoX with structured data, delimiter heads concentrate in late layers (20/40 heads in layers 16-23). On NeoX with chemistry, early layers dominate (~25/50 heads in layers 0-7). The Llama architecture distributes heads more evenly. The model allocates boundary-recovery capacity wherever it is most useful for that domain. The universal property is that delimiter heads form and are causally necessary; the implementation detail of which layers is flexible.

Statistical note. Each PPL ratio represents a single controlled pair (one model with barriers, one without, identical in all other respects). Following standard practice in mechanistic interpretability (Olsson et al., 2022; Voita et al., 2019), where training costs make multi-seed replication impractical, we rely on internal consistency (per-token correlations, cross-format transfer, causal ablation) rather than confidence intervals. The per-token merge-rate correlation in code provides within-experiment statistical structure: the monotonic relationship across five delimiter types would be unlikely under a null hypothesis of random variation.

Tokenizer-attention coupling. We propose this term for the causal relationship demonstrated throughout this paper: tokenizer vocabulary decisions made during tokenizer training propagate through model training and causally constrain which attention heads develop, what they specialize on, and how much capacity is available for content processing. The tokenizer is not a preprocessing step. It is an architectural constraint that shapes the model’s internal organization.

This principle is supported by converging evidence:

- Two architectures (GPT-NeoX with full MHA, Llama with GQA)
- Two model scales (410M, 1.3B)
- Three domains (structured data, programming languages, molecular notation)
- Controlled experiments where the tokenizer is the only variable

- Causal ablation proving the resulting heads are necessary
- Per-token loss correlations showing the effect maps to merge rates at the character level
- Independent validation from Su et al. (2025), whose delimiter performance ranking maps exactly to BPE merge rates
- Unintentional confirmation from Jindal & Ju (2026), whose character-level tokenizer is an implicit merge barrier

The implication for model providers: better tokenizers do not just save tokens. They make models effectively larger by freeing attention capacity from boundary recovery.

12. Discussion

12.1 Why Merging Compounds at Scale

At 10 rows, "name being one token instead of two does not matter. There are only 10 merged boundaries. The attention mechanism can work around it.

At 500 rows, three problems compound simultaneously:

1. The merged boundary repeats 500 times. Each row contains "name":, "id":, "type":. That creates approximately 1,500 positions where the structural boundary is inside a merged token. The model must decompose structure from inside merged tokens at 1,500 positions, not 10.

2. All 1,500 positions are identical token sequences. The token for "name on row 1 is the same integer (#32586) as on row 500. The model cannot distinguish them. It relies on positional encoding alone to track "which "name am I looking at?" Positional encoding degrades over long sequences.

3. 81% of the sequence is noise. The repeated field names and braces are not just merged; they are also redundant. The attention mechanism is spread across approximately 8,500 tokens that carry no information, trying to find the approximately 2,000 tokens that do. The merged boundaries make the noise harder to skip because the model cannot cleanly identify where structure ends and data begins.

Consider the task "how many records have status = shipped?" given 500 JSON objects. The model must attend to every "status": pattern (500 occurrences), read the following value, compare to "shipped," and count matches. The 500 "status": patterns produce the same tokens every time. The model has no structural marker distinguishing the 150th occurrence from the 350th. In a header-factored format, the equivalent task requires attending to a column of values at known, consistent positions. No ambiguity. No repetition competing for attention.

The compounding is critical. At 10 rows: manageable. At 500 rows: 1,500 merged boundaries, massive noise, positional encoding stretched, attention diluted across thousands of identical tokens. This is why JSON errors at scale are off by 50-140 (comprehension failure), not off by 1-2 (precision error).

12.2 Why Claude and Gemma Have Fewer Merges

Claude’s tokenizer has 3 quote+letter entries. Gemma 3 has 2. Across all 43 tokenizers, 11 have zero or near-zero quote merge entries. The specific training details are proprietary, but measurable differences suggest explanations:

- **Vocabulary size** does not predict merge behavior. Claude uses ~65K entries (one of the smallest). Gemma 3 uses 262K (the largest). Both have near-zero quote merges.
- **Training data mix:** Tokenizers trained on corpora with less code/JSON relative to natural language see "name less frequently, making it less likely to cross the merge threshold.
- **Merge boundary policy:** BPE training can be configured to treat certain characters as merge barriers. Anthropic and Google may have intentionally or incidentally prevented " from merging with adjacent letters.

The merge policy matters more than vocabulary size. This observation motivated our controlled experiment: if merge barriers can be applied intentionally and systematically to all 16 delimiter characters, what happens to the trained model?

12.3 The Training Familiarity Paradox

The conventional wisdom holds that LLMs “know” JSON best because they trained on the most JSON. At the model level, this is true. At the tokenizer level, it inverts: the more JSON the tokenizer saw, the more aggressively it merged JSON patterns, and the more boundaries it hid. GPT-4 has 117 merged quote+field entries. Claude has 3. “Trained on JSON” is not an advantage for structural comprehension at scale. It is the mechanism that causes structural ambiguity.

12.4 Why Code Benefits

The code comprehension improvement (3-5x) was not predicted. The barrier characters {, }, (,), :, ; are also code syntax. Every function definition, every code block, every argument list gets clean token boundaries. The model develops the same structural attention heads for code as for structured data. This suggests merge barriers are not a structured-data-specific optimization; they are a general-purpose improvement for any content that relies on delimiter characters.

12.5 TOON’s Tab Delimiter Is Worse Than JSON’s Quote

TOON (Matveev, 2026) uses tab as its column delimiter. Tab has 1,238 mergeable words across all 43 vocabularies, 52x the pipe’s 24. GPT-4o has a 100% tab merge rate. TOON chose the delimiter with the largest adversarial surface of any common separator character. Model A’s 2.3x advantage on tab-separated data (never seen in training) confirms that merge barriers generalize to any delimiter character, including the worst ones.

12.6 The 50x Advantage on Unseen Schemas

The users and logs schemas show 50-51x advantages, far larger than the 2-7x seen on other tests. Model B's PPL on these schemas exceeds 600,000, meaning it essentially cannot parse them at all. These schemas were not in the held-out test data used for the core evaluation (which used product records). They use different field names, different value patterns, and different data shapes. Model A handles them at PPL ~14,000, comparable to its performance on the held-out product schema. This demonstrates that Model A's structural comprehension generalizes across schemas, while Model B's comprehension is fragile and schema-dependent.

12.7 Implications for GQA

GQA's shared KV projections provide partial structural capability to standard-BPE models that separate-KV architectures cannot develop. Production Llama-family models already have some delimiter specialization, which may explain their reasonable structured data performance despite corrupted tokenization. Merge barriers amplify this existing capability.

12.8 What Merge Barriers Cannot Fix

Corruption detection remains limited at this model scale. While delimiter heads contribute to detecting missing fields and wrong headers (Section 7.6), neither model reliably detected all corruption types via PPL spike ($>1.5x$ threshold). The strongest signal was wrong record count (Model A: 1.39x, Model B: 1.37x), below the detection threshold. Neither model produced valid few-shot GCF generations from examples (0/5 both). These tasks require more training capacity than 20,000-40,000 steps on a 410M model. The merge barrier advantage at this scale is in comprehension (reading structure), not in validation (detecting errors) or generation (writing structure from examples). Larger models with more training may close this gap.

12.9 What This Paper Claims and Does Not Claim

Correct: "Merge barriers improve the entire model's structured data processing. The most visible evidence is 50-66 delimiter-specialized attention heads that are causally necessary for format-level comprehension."

Not claimed: "Delimiter heads are the mechanism" (they are the expression). "Delimiter heads are portable modules" (transplant controls disprove this; Appendix D). "Concentration ratio predicts comprehension across models" (confounded; Appendix G).

Also not claimed: "Tokenizer design is the only factor that influences head specialization." Our controlled experiment isolates the tokenizer by holding everything else constant (architecture, corpus, hyperparameters, random seed, hardware). This proves the tokenizer is *a* causal factor, not that it is the only one. Other training conditions likely influence whether and how delimiter heads emerge: training data composition (what fraction is structured data), vocabulary size (which determines which merges occur), training duration (our emergence data shows head counts change with continued

training), and context window length (longer context exposes the model to larger structural payloads). Each of these could be isolated with the same controlled methodology used here. Other architectural choices also changed between our two experiments (positional encoding, activation function, normalization) and could not be isolated from GQA; disentangling these is future work.

12.10 Implications

For tokenizer designers: Merge barriers produce 3-46x better structured data comprehension, better code comprehension, and a model that develops specialized structural attention heads, with natural language perplexity unchanged on held-out Wikipedia (19.4 vs 19.5). The mechanism is architecture-independent.

For model providers: Every model retrain is an opportunity to adopt merge barriers. The tokenizer change is trivial (pre-tokenization rules). The downstream effect (3-46x better structured data and code) is not. As tool use, MCP, and agent pipelines grow, the value of structural comprehension increases.

For format designers: Choose delimiters with the smallest adversarial surface. Pipe (24 words) is 81x safer than JSON's combined grammar (1,939 words). Tab (1,238 words) is the worst common choice. The merge barrier results confirm that delimiter selection matters at the transformer level, not just the tokenizer level.

13. Limitations

1. **Model scale.** Tested at 410M on both architectures. 1.3B-7B planned.
 2. **Training duration.** 20,000 steps on NeoX (~1.3B tokens), 40,000 on Llama. Longer training may change the advantage ratio.
 3. **Context window.** 2,048 tokens. JSON truncates at 50+ records. Production models use 128K+ context.
 4. **Single corpus.** Both architectures trained on the same 4.5 GB corpus. Results may differ with other compositions.
 5. **Flat learning rate.** Used flat LR instead of warmup + cosine decay. Better scheduling might change convergence dynamics.
 6. **High absolute perplexity.** Both models have high PPL on structured data (thousands), reflecting limited training. The relative comparison (3-46x) is what matters, not the absolute numbers.
 7. **PPL-to-comprehension gap.** PPL measured on our models; comprehension on production models. Connection supported by correlation, not direct measurement.
 8. **Head identification sensitivity.** Counts vary with threshold (31-85 on Llama). Causal findings are internally consistent.
 9. **GQA ablation.** Per-query-head intervention is weaker than NeoX's per-head ablation. KV-group ablation partially addresses this.
-

14. Related Work

Head specialization (descriptive). Voita et al. (2019) classified heads as positional/syntactic/rare-word in BERT. Clark et al. (2019) mapped heads to dependency relations. Michel et al. (2019) showed most heads can be pruned. These studies describe what heads do in existing models; none address what training conditions cause heads to specialize.

Mechanistic interpretability (causal). Elhage et al. (2021) introduced a mathematical framework for understanding transformer circuits. Olsson et al. (2022) proved induction heads cause in-context learning. Conmy et al. (2023) developed automated circuit discovery methods. Our work is structurally similar to Olsson et al.: we prove delimiter heads cause structured data comprehension. Their work is within-model (what heads do). Ours is across-models (what makes heads exist).

Tokenizer design. BPE (Sennrich et al., 2016) and its SentencePiece implementation (Kudo and Richardson, 2018) dominate production tokenizers. Alternative approaches include Unigram language models (Kudo, 2018), byte-level models like ByT5 (Xue et al., 2022) that eliminate tokenization entirely, and multi-scale approaches like MegaByte (Yu et al., 2023) that operate on raw bytes with patch-level processing. These alternatives avoid the merge problem by construction but sacrifice the compression efficiency that makes BPE practical for long contexts. Merge barriers achieve clean boundaries while preserving BPE’s compression.

Tokenizer quality and structured data. Deekeswar (2026) measured that 1,000 JSON records consume approximately 80,000 tokens. Our analysis explains the mechanism: 52% are repeated field names that fuse with structural delimiters. Karim and Batatia (2025) proposed fixed tokens for structure and BPE for values. Merge barriers achieve a similar result by construction: structure is always fixed tokens because barrier characters can never merge. Liyanage and Yvon (2026) studied post-training tokenizer adaptation, demonstrating that changes degrade performance. This supports our irrecoverability argument and motivates fixing the tokenizer before training, not after.

Architecture components. Our controlled experiments use two architecture families: GPT-NeoX (Black et al., 2022), representative of the GPT-2/3 lineage (Radford et al., 2019), and Llama (Touvron et al., 2023), which introduced RoPE (Su et al., 2021) and adopted GQA (Ainslie et al., 2023), itself extending Multi-Query Attention (Shazeer, 2019). The architectural differences between these families (learned vs rotary position encodings, full MHA vs GQA, GELU vs SwiGLU) make them a strong test of mechanism generality.

Structured data and LLMs. Sui et al. (2023) showed that table format affects LLM performance. Our analysis explains this at the BPE level and proves the mechanism can be fixed at the tokenizer level. Kutschka and Geiger (2026) found that token-efficient formats can hurt accuracy in some configurations. Our data partially confirms this at small scale but shows the compensation fails at 500+ records.

Frequency-attention dynamics. Ildiz et al. (2024) proved that self-attention weights tokens proportionally to frequency. This is the mathematical basis for grammar attention collapse: when structural tokens dominate by count, they consume the attention budget.

Matveev (2026) argued that JSON’s advantage from training distribution scales with data complexity, proposing that alternative formats only separate past a complexity threshold. Our evaluation data

confirms the threshold exists at approximately 100-200 records for nested data and approximately 500 for flat tables. Our controlled experiment adds a new dimension: even holding the format constant (GCF), the tokenizer determines whether the model can comprehend the structure.

Our contribution bridges two disconnected communities: tokenizer research (compression efficiency) and mechanistic interpretability (attention patterns). Tokenizer design causally determines attention head organization.

15. Conclusion

We have shown that BPE tokenization imposes a permanent constraint on transformer attention capacity that is more severe and more pervasive than previously understood. The constraint is not limited to a small number of specialized heads: all 384 heads at 410M and all 768 heads at 1.3B show 4x more delimiter attention when the tokenizer simply keeps delimiters as separate tokens. The constraint is not a training artifact: it is present from step 5,000 and does not change across 35,000 additional training steps. And the constraint is not inherent to the model: the same weights, with the same vocabulary, produce focused structural attention when given clean token boundaries.

The frustration gap, 41 percentage points of structural attention capacity that the model builds but cannot use, represents a new category of inefficiency in transformer training. Unlike dormant heads (which can be pruned) or undertrained parameters (which improve with more data), stranded capacity is permanently inaccessible under the model's own tokenizer. Training longer does not help. More data does not help. Scaling makes it worse: at 1.3B, standard BPE with GQA develops 124 counterproductive delimiter heads whose removal improves comprehension by 57%.

The mechanism generalizes across domains. On structured data, merge barriers produce 3-738x better perplexity. On code, per-token improvement maps directly to BPE merge rates. On molecular notation, merge barriers produce 67% more valid generated molecules. Effect magnitude correlates with delimiter density across all domains tested. We propose the term tokenizer-attention coupling for this general principle: BPE merge decisions made during tokenizer training causally constrain which attention heads develop and how much capacity is available for content processing.

The fix is merge barriers: forbidding 16 delimiter characters from participating in BPE merges. This is a constraint on the tokenizer training algorithm, not on the model. It requires no architectural changes, no additional parameters, and no modifications to the training pipeline. The disproportion between the severity of the problem and the simplicity of the solution suggests that the tokenizer, the least-studied component of the modern LLM stack, deserves the same empirical scrutiny that architecture, optimization, and data receive.

Future work should validate these findings at 7B+ scale, with production-length context windows (128K+), and with direct comprehension evaluation on downstream tasks. The attention head atlas project (tracking all head specialization types simultaneously from step 0 to convergence) would establish whether tokenizer-attention coupling affects the developmental sequence of all head types, not just delimiter specialists.

Appendix A: Tokenizer Configuration

```
from tokenizers import import Tokenizer, models, trainers, pre_tokenizers
```

```
tokenizer = Tokenizer(models.BPE())
```

```
tokenizer.pre_tokenizer = pre_tokenizers.Sequence([
    pre_tokenizers.Split("|", behavior="isolated"),
    pre_tokenizers.Split("@", behavior="isolated"),
    pre_tokenizers.Split("<", behavior="isolated"),
    pre_tokenizers.Split(">", behavior="isolated"),
    pre_tokenizers.Split("'", behavior="isolated"),
    pre_tokenizers.Split('"', behavior="isolated"),
    pre_tokenizers.Split(":", behavior="isolated"),
    pre_tokenizers.Split(",", behavior="isolated"),
    pre_tokenizers.Split(";", behavior="isolated"),
    pre_tokenizers.Split("\t", behavior="isolated"),
    pre_tokenizers.Split("{", behavior="isolated"),
    pre_tokenizers.Split("}", behavior="isolated"),
    pre_tokenizers.Split("[", behavior="isolated"),
    pre_tokenizers.Split("]", behavior="isolated"),
    pre_tokenizers.Split("(", behavior="isolated"),
    pre_tokenizers.Split(")", behavior="isolated"),
    pre_tokenizers.ByteLevel(add_prefix_space=False),
])
```

```
trainer = trainers.BpeTrainer(
    vocab_size=65536,
    special_tokens=["<pad>", "<eos>"],
)
tokenizer.train(files=["corpus.txt"], trainer=trainer)
```

Each `Split` isolates one barrier character before BPE merging begins. The `ByteLevel` pre-tokenizer handles the remaining text using standard GPT-style byte encoding. No changes to the BPE algorithm, training pipeline, or model architecture are required. The improvement is entirely in the pre-tokenization configuration.

Appendix B: Per-Tokenizer Vocabulary Counts

Merged vocabulary entries by tokenizer family:

Tokenizer	Vocab Size	Quote+Letter	Pipe+Letter	Tab+Letter	Multi-Grammar
GPT-4 (cl100k)	~100K	117	22	1,173	874
GPT-4o (o200k)	~199K	108	8	1,036	735
Claude	~65K	3	0	0	667
LLaMA 2	~32K	0	0	0	122
LLaMA 3/3.1	~128K	153	277	0	881
Qwen 2.5	~152K	154	40	0	874
DeepSeek V3/R1	~129K	48	5	0	287
Gemma 2	~256K	4	0	0	735
Gemma 3	~262K	2	0	0	861
Mistral Nemo	~131K	38	3	0	415
Phi-4	~100K	153	116	0	874
Falcon	~65K	4	1	0	92
StarCoder2	~49K	3	2	0	535
Jamba	~66K	1	1,543	0	113

Claude and Gemma have near-zero quote merge entries (3 and 2-4 respectively), suggesting intentional or incidental merge boundary policies during tokenizer training. This is evidence that clean boundaries are achievable without sacrificing general-purpose quality.

Appendix C: Structural Pattern Test

We tested whether cross-format transfer depends on the delimiter character or the structural pattern by holding one constant and varying the other. 5 formats, same data, 30 records each:

Format	Character	Pattern	Delta	Transfers?
A: tab + GCF layout	tab	flat separator	+51.2%	YES
B: tab + TSV layout	tab	header+rows	+32.1%	YES
C: tab + wrapping	tab	wrapping	+123.4%	YES

Format	Character	Pattern	Delta	Transfers?
D: pipe + wrapping	pipe	wrapping	-54.0%	NO (adversarial)
E: GCF (control)	pipe	flat separator	+59.9%	YES

Pipe, GCF's own delimiter, becomes actively adversarial (-54%) in a wrapping layout. The heads learned "pipe means flat field separator." When pipe appears in wrapping context, they misapply that reasoning. Tab transfers in all contexts because no conflicting prior exists. The mechanism: familiar delimiters become adversarial in unfamiliar contexts; unfamiliar delimiters get neutral-to-positive treatment.

Appendix D: Transplant Controls

Grafted Model A's delimiter head weights (Q, K, V, output projections) into Model B's corresponding positions. No retraining.

Progressive transplant (NeoX)

Heads transplanted	GCF delta	JSON delta	NL delta
5	-53%	-37%	+5%
10	-77%	-69%	+5%
20	-83%	-84%	+11%
40	-96%	-94%	+16%
101 (all)	-99.3%	-100%	+68%

Controls at 20 heads

Control	GCF	JSON	TOON	CSV	NL
Delimiter heads A->B	-81%	-86%	-33%	-59%	+12%
Random heads A->B	-70%	-99%	-87%	-95%	+1%
B's heads -> A	-43%	n/t	n/t	n/t	n/t
Shifted positions	-94%	n/t	n/t	n/t	n/t

Critical finding: Random non-delimiter heads from Model A also substantially improve Model B (-70% GCF vs -81% for delimiter heads). On JSON and TOON, random heads improve B *more* than delimiter heads. The improvement is holistic: merge barriers improve all of Model A’s weights through training, not just the delimiter-specialized heads. Cross-position transplant works equally well (-94%), confirming the weights are not position-dependent. The ablation (within-model, causal) remains the stronger evidence for delimiter head causality.

Appendix E: Tokenization Examples

E.1 Edge declaration: @0<@2|implements

Tokenizer	Tokens	Split
Claude	7	@ 0 < @ 2 \ implements
GPT-4	7	@ 0 < @ 2 \ implements
GPT-4o	7	@ 0 < @ 2 \ implements
LLaMA 3.1	7	@ 0 < @ 2 \ implements
Qwen 2.5	7	@ 0 < @ 2 \ implements
DeepSeek V3	8	@ 0 < @ 2 \ implements
Gemma 2	7	@ 0 < @ 2 \ implements
Mistral Nemo	8	@ 0 < @ 2 \ implements

All structural characters (@, <, |) are always single tokens. The only variance is in the value implements (1 vs 2 tokens), which does not affect parsing.

E.2 Symbol row: @0|function|auth.validateToken|0.95|definition

Tokenizer	Tokens	Key Differences
GPT-4	14	Merges .validate (1 tok), 95 (1 tok)
Qwen 2.5	15	Splits 95 into 9 + 5
Gemma 2	16	Splits . + validate, splits 9 + 5

Pipe delimiters are always single tokens across all tokenizers. Variance is only in how tokenizers handle value content: dot-prefixed words and two-digit numbers. This is value variance (harmless), not boundary variance (dangerous).

E.3 Delimiter selection rationale

Character	Why chosen	Alternative considered	Why not
\ (pipe)	24-word surface, all TypeScript union keywords. Visually distinct column separator.	Backtick (5 words), Tilde (8 words)	Backtick conflicts with mark-down/template literals, tilde with paths.
@	“This is an ID” semantics. 127-word surface, but used only before digits (@0, @1), which never trigger merges.	\$	Also safe, but less intuitive.
##	Two-char sequence always merges into one token. Markdown-familiar.	===	3 chars, less efficient.
<	Reads as “points to” for edges.	~	Also safe, but less semantic.
\n	Universal row separator, zero overhead.	;	Less readable.
,	Schema field separator, familiar from CSV.	:	Conflicts with value content.

Appendix F: Reproducibility

Analysis Scripts (Open Source)

Script	Purpose
hf-tokenizer-analysis.py	43-tokenizer merge rates, vocab entries
structural-equivalence-proof.py	Grammar isolation across 43 tokenizers

Script	Purpose
<code>adversarial-vocab-dump.py</code>	Exhaustive vocabulary scan, adversarial surface
<code>attention-analysis.py</code>	Attention extraction from Pythia 410M / Gemma 2B
<code>ascii-adversarial-surface.py</code>	All 94 printable ASCII characters ranked

Controlled Experiments

Component	Detail
Architectures	GPT-NeoX 410M (run-002), Llama 410M (run-003)
Model architecture	NeoX: 24 layers, 16 heads, 1024 hidden. Llama: 24 layers, 16 heads (GQA 4:1), 1024 hidden
Framework	PyTorch 2.4.1, HuggingFace transformers
Hardware	4x A100 PCIE 40GB (run-002), RTX 6000 Ada + RTX 4090 (run-003)
Training	DDP with NCCL, gradient checkpointing, fp16
Training steps	20K NeoX, 40K Llama, same corpus
Total GPU hours	~35 hours (training) + ~6 hours (eval/ablation)
Ablation phases	18 (run-002) + 12 (run-003)
Eval scripts	22, all in repository
Result files	51 (run-002) + 26 (run-003) with full provenance
Win/loss record (NeoX)	11/11 (structured + code), 0/0 natural language (tied)
Checkpoints	HuggingFace (blackwell-systems/structok-checkpoints, private)
Logs	Cloudflare R2 (structok-training bucket)

Repository: github.com/blackwell-systems/merge-barriers

Appendix G: Production Model Probing

Probed production models for delimiter head specialization using excess delimiter attention and concentration ratio (fraction of total excess in the top 10% of heads).

Model	Params	Heads	Top-10 excess	Concentration	GCF score
Model A (merge barriers)	410M	384	0.349	54.3%	N/A (PPL)
Model B (standard BPE)	410M	384	0.282	63.6%	N/A (PPL)
Phi-2	2.7B	1024	0.626	17.9%	N/A
Gemma 2 2B	2.6B	208	0.662	18.0%	N/A
Llama 3.1 8B	8B	1024	0.755	14.9%	65.4%
Mistral 7B	7B	1024	0.836	14.5%	64.6%
Qwen 2.5 7B	7B	784	0.247	72.6%	61.5%

Merge barriers create a small number of deeply specialized heads (54% concentration). Standard BPE in production models creates many heads that all attend to delimiters somewhat but none deeply (14-15% concentration). Qwen’s high concentration (72.6%) but low comprehension (61.5%) broke the “concentration predicts comprehension” hypothesis.

Fundamental confound: Each model has a different tokenizer, so “delimiter token” means different things. On Model A, " is always its own token. On Mistral, "name is one token. The probing measures different things on different models. Present as exploratory observation, not predictive metric. The B0 finding (Section 8.4, 35 functional heads on standard Llama trained with the same tokenizer) supersedes this as evidence for how standard-BPE models develop structural capability.

References

- Ainslie, J., Lee-Thorp, J., de Jong, M., Zemlyanskiy, Y., Lebron, F., & Sanghai, S. (2023). GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints. EMNLP. arXiv:2305.13245.
- Black, S., Biderman, S., Hallahan, E., Anthony, Q., Gao, L., Golding, L., He, H., Leahy, C., McDonell, K., Phang, J., Pieler, M., Prashanth, U. S., Purohit, S., Reynolds, L., Tow, J., Wang, B., & Weinbach, S. (2022). GPT-NeoX-20B: An Open-Source Autoregressive Language Model. ACL Workshop on Challenges & Perspectives in Creating Large Language Models. arXiv:2204.06745.
- Blackwell, D. (2026). GCF: A Token-Optimized Wire Format for Structured LLM Interactions. DOI: [10.5281/zenodo.20579817](https://doi.org/10.5281/zenodo.20579817).
- Chen, N., Sun, S., & Hooi, B. (2026). Mechanistic Data Attribution: Tracing the Training Origins of Interpretable LLM Units. arXiv:2601.21996.
- Clark, K., Khandelwal, U., Levy, O., & Manning, C. D. (2019). What Does BERT Look At? An Analysis of BERT's Attention. ACL Workshop BlackboxNLP. arXiv:1906.04341.
- Conmy, A., Mavor-Parker, A., Lynch, A., Heimersheim, S., & Garriga-Alonso, A. (2023). Towards Automated Circuit Discovery for Mechanistic Interpretability. NeurIPS. arXiv:2304.14997.
- Deekeswar, H. (2026). ONTO: A Token-Efficient Columnar Notation for LLM Input Optimization. arXiv:2604.17512.
- Elhage, N., Nanda, N., Olsson, C., Henighan, T., Joseph, N., Mann, B., Askell, A., Bai, Y., Chen, A., Conerly, T., DasSarma, N., Drain, D., Ganguli, D., Hatfield-Dodds, Z., Hernandez, D., Jones, A., Kernion, J., Lovitt, L., Ndousse, K., Amodei, D., Brown, T., Clark, J., Kaplan, J., McCandlish, S., & Olah, C. (2021). A Mathematical Framework for Transformer Circuits. Transformer Circuits Thread, Anthropic. <https://transformer-circuits.pub/2021/framework/index.html>.
- Gu, X., Liu, Y., & Feng, Y. (2025). When Attention Sink Emerges in Language Models: An Empirical View. ICLR (Spotlight). arXiv:2410.10781.
- Hasani, M., Fan, Z., Shen, C., & Ren, K. (2026). Mechanistic Interpretability of Large-Scale Counting in LLMs through a System-2 Strategy. arXiv:2601.02989.
- Ildiz, M. E., Huang, Y., Li, Y., Rawat, A. S., & Oymak, S. (2024). From Self-Attention to Markov Models: Unveiling the Dynamics of Generative Transformers. arXiv:2402.13512.
- Jindal, S. & Ju, C. (2026). SMoLLM: Small Language Models Learn Small Molecular Grammar. arXiv:2605.06322.
- Karim, K. & Batatia, H. (2025). Innovative Tokenisation of Structured Data for LLM Training. arXiv:2508.01685.
- Katzy, J., Klinkenberg, R., & Bielak, P. (2026). Automated Attention Pattern Discovery at Scale in Large Language Models. arXiv:2604.03764.
- Kudo, T. (2018). Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates. ACL. arXiv:1804.10959.

- Kudo, T. & Richardson, J. (2018). SentencePiece: A Simple and Language Independent Subword Tokenizer and Detokenizer for Neural Text Processing. EMNLP System Demonstrations. arXiv:1808.06226.
- Kutschka, L. & Geiger, B. (2026). Notation Matters: A Benchmark Study of Token-Optimized Formats in Agentic AI Systems. arXiv:2605.29676.
- Liyanage, V. & Yvon, F. (2026). AdaptBPE: From General Purpose to Specialized Tokenizers. arXiv:2601.21665.
- Matveev, I. (2026). Token-Oriented Object Notation vs JSON: A Benchmark of Plain and Constrained Decoding Generation. arXiv:2603.03306.
- Michel, P., Levy, O., & Neubig, G. (2019). Are Sixteen Heads Really Better than One? NeurIPS. arXiv:1905.10650.
- Pande, A., Budhraja, A., & Nema, P. (2021). The Heads Hypothesis: A Unifying Statistical Approach Towards Understanding Multi-Headed Attention in BERT. AAAI. arXiv:2101.09115.
- Olsson, C., Elhage, N., Nanda, N., Joseph, N., DasSarma, N., Henighan, T., Mann, B., Askell, A., Bai, Y., Chen, A., Conerly, T., Drain, D., Ganguli, D., Hatfield-Dodds, Z., Hernandez, D., Johnston, S., Jones, A., Kernion, J., Lovitt, L., Ndousse, K., Amodei, D., Brown, T., Clark, J., Kaplan, J., McCandlish, S., & Olah, C. (2022). In-context Learning and Induction Heads. arXiv:2209.11895.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language Models Are Unsupervised Multitask Learners. OpenAI.
- Sandoval-Segura, P., Wang, X., Panda, A., Goldblum, M., Basri, R., Goldstein, T., & Jacobs, D. (2025). Using Attention Sinks to Identify and Evaluate Dormant Heads in Pretrained LLMs. arXiv:2504.03889.
- Sennrich, R., Haddow, B., & Birch, A. (2016). Neural Machine Translation of Rare Words with Subword Units. In Proceedings of the 54th Annual Meeting of the ACL (pp. 1715-1725). arXiv:1508.07909.
- Sharma, A., Hewitt, J., & Friedman, S. (2026). Dissociating Decodability and Causal Use in Bracket-Sequence Transformers. arXiv:2604.22128.
- Su, J., Zhang, J., Ullrich, K., Bottou, L., & Ibrahim, M. (2025). A Single Character can Make or Break Your LLM Evals. arXiv:2510.05152.
- Shazeer, N. (2019). Fast Transformer Decoding: One Write-Head is All You Need. arXiv:1911.02150.
- Su, J., Lu, Y., Pan, S., Murtadha, A., Wen, B., & Liu, Y. (2021). RoFormer: Enhanced Transformer with Rotary Position Embedding. arXiv:2104.09864.
- Sui, Y., He, M., Zhang, Z., Wang, Y., & Zhao, J. (2023). Table Meets LLM: Can Large Language Models Understand Structured Table Data? A Benchmark and Empirical Study. arXiv:2305.13062.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Roziere, B., Goyal, N., Hambro, E., Azhar, F., Rodriguez, A., Joulin, A., Grave, E., & Lample, G. (2023). LLaMA: Open and Efficient Foundation Language Models. arXiv:2302.13971.
- University of Mannheim. (2024). Web Data Commons: RDFa, Microdata, and Microformat Data Sets. <http://webdatacommons.org/structureddata/>

- Voita, E., Talbot, D., Moiseev, F., Sennrich, R., & Titov, I. (2019). Analyzing Multi-Head Self-Attention: Specialized Heads Do the Heavy Lifting, and the Rest Can Be Pruned. ACL. arXiv:1905.09418.
- Wan, Z., Zhao, H., Yang, Y., Wang, Z., & Luu, A. T. (2022). What Do They Capture? A Structural Analysis of Pre-Trained Language Models for Source Code. ICSE. arXiv:2202.06840.
- Wang, K., Variengien, A., Conmy, A., Shlegeris, B., & Steinhardt, J. (2022). Interpretability in the Wild: a Circuit for Indirect Object Identification in GPT-2 Small. arXiv:2211.00593.
- Wang, L., Chen, Z., & Liu, Y. (2025). Differentiation and Specialization of Attention Heads via the Refined Local Learning Coefficient. ICLR. arXiv:2410.02984.
- Wang, Y., Li, Z., & Zhang, Y. (2026). A Closer Look into LLMs for Table Understanding. arXiv:2603.15402.
- Xiao, G., Tian, Y., Chen, B., Han, S., & Lewis, M. (2024). Efficient Streaming Language Models with Attention Sinks. ICLR. arXiv:2309.17453.
- Yao, S., Chen, B., Levy, O., & Yin, W. (2021). Self-Attention Networks Can Process Bounded Hierarchical Languages. ACL. arXiv:2105.11115.
- Zhang, H., Wang, Y., & Li, S. (2026). How Do Language Models Understand Tables? A Mechanistic Analysis of Cell Location. arXiv:2602.08548.
- Xue, L., Barua, A., Constant, N., Al-Rfou, R., Narang, S., Kale, M., Roberts, A., & Raffel, C. (2022). ByT5: Towards a Token-Free Future with Pre-trained Byte-to-Byte Models. TACL. arXiv:2105.13626.
- Yu, L., Simig, D., Flaherty, C., Aghajanyan, A., Zettlemoyer, L., & Lewis, M. (2023). MEGABYTE: Predicting Million-byte Sequences with Multiscale Transformers. NeurIPS. arXiv:2305.07185.

Corresponding author: Dayna Blackwell, Blackwell Systems (dayna@blackwell-systems.com)